

# CSCI 2021, Sprint 2015, Homework Assignment III

**Problem 0:**

Clearly label your assignment with the time of your recitation section (8:00, 9:05, 10:10, 11:15, 12:20, 1:25, 2:30). This will help us turn back your graded assignments more efficiently.

The next problem concerns the following C code. This program reads a string on standard input and prints an integer in hexadecimal format based on the input string it read.

```
#include <stdio.h>

/* Read a string from stdin into buf */
int evil_read_string()
{
    int buf[2];

    scanf("%s", (char *)buf);
    return buf[1];
}

int main()
{
    printf("0x%x\n", evil_read_string());
}
```

Here is the corresponding machine code on a Linux/x86 machine:

```
08048414 <evil_read_string>:
8048414: 55          push    %ebp
8048415: 89 e5       mov     %esp,%ebp
8048417: 83 ec 14    sub     $0x14,%esp
804841a: 53          push    %ebx
804841b: 83 c4 f8    add     $0xffffffff8,%esp
804841e: 8d 5d f8    lea     0xffffffff8(%ebp),%ebx
8048421: 53          push    %ebx                address arg for scanf
8048422: 68 b8 84 04 08 push    $0x80484b8         format string for scanf
8048427: e8 e0 fe ff ff call     804830c <_init+0x50> call scanf
804842c: 8b 43 04    mov     0x4(%ebx),%eax
804842f: 8b 5d e8    mov     0xffffffffe8(%ebp),%ebx
8048432: 89 ec       mov     %ebp,%esp
8048434: 5d          pop     %ebp
8048435: c3          ret

08048438 <main>:
8048438: 55          push    %ebp
8048439: 89 e5       mov     %esp,%ebp
804843b: 83 ec 08    sub     $0x8,%esp
804843e: 83 c4 f8    add     $0xffffffff8,%esp
8048441: e8 ce ff ff ff call     8048414 <evil_read_string>
8048446: 50          push    %eax                integer arg for printf
8048447: 68 bb 84 04 08 push    $0x80484bb         format string for printf
804844c: e8 eb fe ff ff call     804833c <_init+0x80> call printf
8048451: 89 ec       mov     %ebp,%esp
8048453: 5d          pop     %ebp
8048454: c3          ret
```

## Problem 1:

This problem tests your understanding of the stack discipline and byte ordering. Here are some notes to help you work the problem:

- `scanf("%s", buf)` reads an input string from the standard input stream (stdin) and stores it at address `buf` (including the terminating `'\0'` character). It does **not** check the size of the destination buffer.
- `printf("0x%x", i)` prints the integer `i` in hexadecimal format preceded by "0x".
- Recall that Linux/x86 machines are little endian.
- You will need to know the hex values of the following characters:

Character	Hex value	Character	Hex value
'd'	0x64	'v'	0x76
'r'	0x72	'i'	0x69
'.'	0x2e	'l'	0x6c
'e'	0x65	'\0'	0x00
		's'	0x73

- A. Suppose we run this program on a Linux/x86 machine, and give it the string "dr.evil" as input on stdin.

Here is a template for the stack, showing the locations of `buf[0]` and `buf[1]`. Fill in the value of `buf[1]` (in hexadecimal) and indicate where `ebp` points **just after** `scanf` returns to `evil_read_string`.

```

|<- buf[0]->|<-buf[1] ->|
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
|   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
```

What is the 4-byte integer (in hex) printed by the `printf` inside `main`?

0x\_\_\_\_\_

B. Suppose now we give it the input “dr.evil.lives” (again on a Linux/x86 machine).

- (a) List the contents of the following memory locations just **after** `scanf` returns to `evil_read_string`. Each answer should be an unsigned 4-byte integer expressed as 8 hex digits.

`buf[0]` = 0x\_\_\_\_\_

`buf[3]` = 0x\_\_\_\_\_

- (b) Immediately **before** the `ret` instruction at address `0x08048435` executes, what is the value of the frame pointer register `%ebp`?

`%ebp` = 0x\_\_\_\_\_

You can use the following template of the stack as *scratch space*. *Note:* this does **not** have to be filled out to receive full credit.

```

      <- buf[0] -><- buf[1] ->
--+-+-+-+-+
| | | | | | | | | | | | | | | | | | | | | | | | | | | |
--+-+-+-+-+
```

## Buffer overflow

The next problem concerns the following C code, excerpted from Dr. Evil's best-selling autobiography, "World Domination My Way". He calls the program *NukeJr*, his baby nuclear bomb phase.

```
/*
 * NukeJr - Dr. Evil's baby nuke
 */
#include <stdio.h>

int overflow(void);
int one = 1;

/* main - NukeJr's main routine */
int main() {
    int val = overflow();

    val += one;
    if (val != 15213)
        printf("Boom!\n");
    else
        printf("Curses! You've defused NukeJr!\n");
    _exit(0); /* syscall version of exit that doesn't need %ebp */
}

/* overflow - writes to stack buffer and returns 15213 */
int overflow() {
    char buf[4];
    int val, i=0;

    while(scanf("%x", &val) != EOF)
        buf[i++] = (char)val;
    return 15213;
}
```

## Buffer overflow (cont)

Here is the corresponding machine code for NukeJr when compiled and linked on a Linux/x86 machine:

```
08048560 <main>:
8048560:    55                pushl   %ebp
8048561:    89 e5             movl    %esp,%ebp
8048563:    83 ec 08          subl    $0x8,%esp
8048566:    e8 31 00 00 00    call    804859c <overflow>
804856b:    03 05 90 96 04    addl    0x8049690,%eax        # val += one;
8048570:    08
8048571:    3d 6d 3b 00 00    cmpl    $0x3b6d,%eax        # val == 15213?
8048576:    74 0a             je      8048582 <main+0x22>
8048578:    83 c4 f4          addl    $0xffffffff4,%esp
804857b:    68 40 86 04 08    pushl   $0x8048640
8048580:    eb 08            jmp     804858a <main+0x2a>
8048582:    83 c4 f4          addl    $0xffffffff4,%esp
8048585:    68 60 86 04 08    pushl   $0x8048660
804858a:    e8 75 fe ff ff    call    8048404 <_init+0x44> # call printf
804858f:    83 c4 10          addl    $0x10,%esp
8048592:    83 c4 f4          addl    $0xffffffff4,%esp
8048595:    6a 00            pushl   $0x0
8048597:    e8 b8 fe ff ff    call    8048454 <_init+0x94> # call _exit

0804859c <overflow>:
804859c:    55                pushl   %ebp
804859d:    89 e5             movl    %esp,%ebp
804859f:    83 ec 10          subl    $0x10,%esp
80485a2:    56                pushl   %esi
80485a3:    53                pushl   %ebx
80485a4:    31 f6            xorl    %esi,%esi
80485a6:    8d 5d f8          leal    0xffffffff8(%ebp),%ebx
80485a9:    eb 0d            jmp     80485b8 <overflow+0x1c>
80485ab:    90                nop
80485ac:    8d 74 26 00       leal    0x0(%esi,1),%esi
80485b0:    8a 45 f8          movb    0xffffffff8(%ebp),%al    # L1: loop start
80485b3:    88 44 2e fc       movb    %al,0xffffffffc(%esi,%ebp,1)
80485b7:    46                incl    %esi
80485b8:    83 c4 f8          addl    $0xffffffff8,%esp
80485bb:    53                pushl   %ebx
80485bc:    68 80 86 04 08    pushl   $0x8048680
80485c1:    e8 6e fe ff ff    call    8048434 <_init+0x74>    # call scanf
80485c6:    83 c4 10          addl    $0x10,%esp
80485c9:    83 f8 ff          cmpl    $0xffffffff,%eax
80485cc:    75 e2            jne     80485b0 <overflow+0x14> # goto L1
80485ce:    b8 6d 3b 00 00    movl    $0x3b6d,%eax
80485d3:    8d 65 e8          leal    0xffffffe8(%ebp),%esp
80485d6:    5b                popl    %ebx
80485d7:    5e                popl    %esi
80485d8:    89 ec            movl    %ebp,%esp
80485da:    5d                popl    %ebp
80485db:    c3                ret
```

## Problem 2:

This problem uses the NukeJr program to test your understanding of the stack discipline and byte ordering. Here are some notes to help you work the problem:

- Recall that Linux/x86 machines are Little Endian.
- The `scanf("%x", &val)` function reads a whitespace-delimited sequence of characters from `stdin` that represents a hex integer, converts the sequence to a 32-bit `int`, and assigns the result to `val`. The call to `scanf` returns either 1 (if it converted a sequence) or EOF (if no more sequences on `stdin`).

For example, calling `scanf` four times on the input string "0 a ff" would have the following result:

- 1st call to `scanf`: `val=0x0` and `scanf` returns 1.
- 2nd call to `scanf`: `val=0xa` and `scanf` returns 1.
- 3rd call to `scanf`: `val=0xff` and `scanf` returns 1.
- 4th call to `scanf`: `val` is unchanged and `scanf` returns EOF.

- A. After the `subl` instruction at address `0x804859f` in function `overflow` completes, the stack contains a number of objects which are shown in the table below. Determine the address of each object as a byte offset from `buf[0]`.

Stack object	Address of stack object
return address	&buf[0] + _____
old %ebp	&buf[0] + _____
buf[3]	&buf[0] + _____
buf[2]	&buf[0] + _____
buf[1]	&buf[0] + 1
buf[0]	&buf[0] + 0

- B. What input string would defuse NukeJr by causing the call to `overflow` to return to address `0x8048571` instead of `804856b`? *Notes: (i) Your solution is allowed to trash the contents of the %ebp register. (ii) Fill in each blank with a one or two digit hex number.*

Answer: "0 0 0 0 \_\_\_\_\_ " \_\_\_\_\_ "

### Problem 3:

Fill in the following tables. Each table describes the actions that are taken at each stage of the execution for an instruction. Your answer should follow the conventions of section 4.3 in the textbook.

	leave Described in problem 4.48 in the textbook
Fetch	
Decode	
Execute	
Memory	
Write back	
PC Update	

	iaddl V, rB Described in problem 4.47 in the textbook
Fetch	
Decode	
Execute	
Memory	
Write back	
PC Update	



	<code>jmp *D(rA)</code> An indirect jump similar to what is described in Section 3.6.3. The jump destination is stored in memory location <code>rA + D</code> .
Fetch	
Decode	
Execute	
Memory	
Write back	
PC Update	

The following code segment correspond to the next two questions.

L5:

<code>mrmovl (%ebx), %eax</code>	(In 1)
<code>rmmovl %eax, (%ecx)</code>	(In 2)
<code>irmovl \$4, %edi</code>	(In 3)
<code>subl %edi, %edx</code>	(In 4)
<code>addl %edi, %ecx</code>	(In 5)
<code>xorl %eax, %edi</code>	(In 6)
<code>andl %edx, %edx</code>	(In 7)
<code>jg L5</code>	(In 8)

## Problem 4:

Consider a low-cost pipelined processor based on the structure shown in Figure 4.41, with no bypassing paths. In other words, the processor in this problem handles all data hazards by stalling. The branch predictor always predicts that a branch is *not* taken. Fill in the following tables with which execution stage each instruction is in, for each clock cycle. The columns correspond to clock cycles. The rows correspond to individual instructions. Instructions can appear multiple time in the table, because these instructions are part of a loop. For example, the existing entries in the table indicate that `In 1` is in fetch stage in cycle 1, in decode stage in cycle 2, in execution stage in cycle 3 in memory stage in cycle 4 and in writeback stage in cycle 5.

Iteration 1																				
Instr	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20
Iteration 1																				
In 1	F	D	E	M	W															
In 2																				
In 3																				
In 4																				
In 5																				
In 6																				
In 7																				
In 8																				
Iteration 2																				
In 1																				
In 2																				
In 3																				
In 4																				
In 5																				
In 6																				
In 7																				
In 8																				
Iteration 3																				
In 1																				
In 2																				
In 3																				
In 4																				
In 5																				
In 6																				
In 7																				
In 8																				

The throughput of a pipeline is typically measure in terms of cycle per instruction (CPI). CPI is the average number of instructions that complete execution per clock cycle. The performance of a pipeline is function of the CPI and the clock rate.

1. The effective throughput of this pipeline in executing this code segment = \_\_\_\_\_;
2. If the clock operates at 1GHz, the instruction-per-second rate achieved by the pipeline on this code = \_\_\_\_\_.

## Problem 5:

Consider a pipelined processor with bypassing paths as shown in Figure 4.52 and a perfect branch predictor. Fill in the following tables with which execution stage each instruction is in, during each clock cycle. The columns correspond to clock cycles. The rows correspond to individual instructions. Instructions can appear multiple time in the table, because these instructions are part of a loop. For example, the existing entries in the table indicate that *In 1* is in fetch stage in cycle 1, in decode stage in cycle 2, in execution stage in cycle 3 in memory stage in cycle 4 and in writeback stage in cycle 5. Marks all the bypassings that occur during the execution.

Iteration 1																				
Instr	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20
Iteration 1																				
In 1	F	D	E	M	W															
In 2																				
In 3																				
In 4																				
In 5																				
In 6																				
In 7																				
In 8																				
Iteration 2																				
In 1																				
In 2																				
In 3																				
In 4																				
In 5																				
In 6																				
In 7																				
In 8																				
Iteration 3																				
In 1																				
In 2																				
In 3																				
In 4																				
In 5																				
In 6																				
In 7																				
In 8																				

The throughput of a pipeline is typically measure in terms of cycle per instruction (CPI). CPI is the average number of instructions graduated per clock cycle. The performance of a pipeline is function of the CPI and the clock rate.

1. The effective throughput of this pipeline in executing this code segment = \_\_\_\_\_;
2. If the clock operates at 1GHz, the instruction per second achieved by the pipeline on this code = \_\_\_\_\_.

**Problem 6:**

Consider a pipeline with a complex ALU that is able to support multiplication and divide instructions. Multiplication instructions take 4 cycles to execute; and divide instructions take 6 clock cycles to complete. The ALU generates an ExCmp signal. ExCmp is set at the end of the final cycle of a long computation. For instance if a multiplication starts on cycle 0, ExCmp is set by the end of cycle 3. All other ALU operations still complete in only one cycle as before, and so set ExCmp to 1 by the end of the first cycle when the inputs are presented. Describe changes necessary to the pipelining control logic to incorporate such an ALU.

**Problems 0, 3, and 5 should be submitted for grading.**