

Linking

CSCI 2021: Machine Architecture and Organization

Antonia Zhai

Department Computer Science and Engineering

University of Minnesota

<http://www.cs.umn.edu/~zhai>

With Slides from Bryant and O'Hallaron



Linker Puzzles

```
int x;  
p1() {}  
p1() {}
```

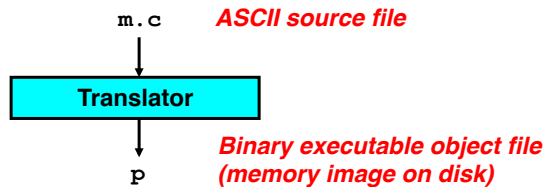
```
int x;  
p1() {}  
int x;  
p2() {}
```

```
int x;  
int y;  
p1() {}  
double x;  
p2() {}
```

```
int x=7;  
int y=5;  
p1() {}  
double x;  
p2() {}
```

```
int x=7;  
p1() {}  
int x;  
p2() {}
```

A Simplistic Program Translation Scheme



Problems:

- Efficiency: small change requires complete recompilation
- Modularity: hard to share common functions (e.g. `printf`)

Solution:

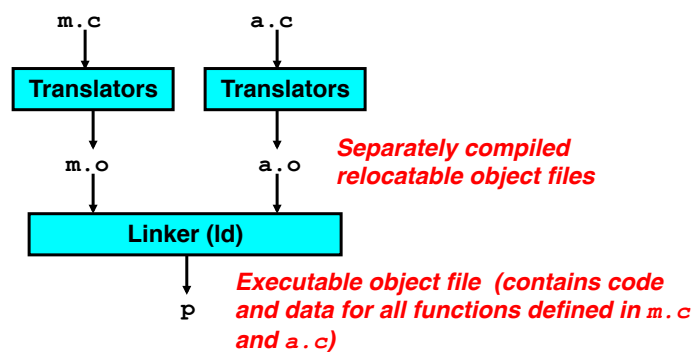
- *Static linker (or linker)*

5/8/15

CSCI 2021

3

A Better Scheme Using a Linker



5/8/15

CSCI 2021

4

Translating the Example Program

Compiler driver coordinates all steps in the translation and linking process.

- Typically included with each compilation system (e.g., `gcc`)
- Invokes preprocessor (`cpp`), compiler (`cc1`), assembler (`as`), and linker (`ld`).
- Passes command line arguments to appropriate phases

Example: create executable `p` from `m.c` and `a.o`:

5/8/15

CSCI 2021

5

Why Linkers?

Modularity

- Program can be written as a collection of smaller source files, rather than one monolithic mass.
- Can build libraries of common functions (more on this later)
 - e.g., Math library, standard C library

Efficiency

- Time:
 - Change one source file, compile, and then relink.
 - No need to recompile other source files.
- Space:
 - Libraries of common functions can be aggregated into a single file...
 - Yet executable files and running memory images contain only code for the functions they actually use.


5/8/15

CSCI 2021

6


What Do Linkers Do?

Step 1. Symbol resolution

- Programs define and reference *symbols* (variables and functions):
 - `void swap() {...}` `/* define symbol swap */`
 - `swap();` `/* reference symbol a */`
 - `int *xp = &x;` `/* define symbol xp, reference x */`
 - Symbol definitions are stored (by compiler) in *symbol table*.
 - Symbol table is an array of structs
 - Each entry includes name, size, and location of symbol.
 - Linker associates each symbol reference with exactly one symbol definition.
- 

What Do Linkers Do? (cont)

Step 2. Relocation

- Merges separate code and data sections into single sections
 - Relocates symbols from their relative locations in the `.o` files to their final absolute memory locations in the executable.
 - Updates all references to these symbols to reflect their new positions.
- 

Three Kinds of Object Files (Modules)

- Relocatable object file (`.o` file)
 - Contains code and data in a form that can be combined with other relocatable object files to form executable object file.
 - Each `.o` file is produced from exactly one source (`.c`) file
- Executable object file (`a.out` file)
 - Contains code and data in a form that can be copied directly into memory and then executed.
- Shared object file (`.so` file)
 - Special type of relocatable object file that can be loaded into memory and linked dynamically, at either load time or run-time.
 - Called *Dynamic Link Libraries* (DLLs) by Windows

Executable and Linkable Format (ELF)

- Standard binary format for object files
- Derives from AT&T System V Unix
 - Later adopted by BSD Unix variants and Linux
- One unified format for
 - Relocatable object files (`.o`),
 - Executable object files
 - Shared object files (`.so`)
- Generic name: ELF binaries

ELF Object File Format

- Elf header
 - Word size, byte ordering, file type (.o, exec, .so), machine type, etc.
- Segment header table
 - Page size, virtual addresses memory segments (sections), segment sizes.
- .text section
 - Code
- .rodata section
 - Read only data: jump tables, ...
- .data section
 - Initialized global variables
- .bss section
 - Uninitialized global variables
 - "Block Started by Symbol"
 - "Better Save Space"
 - Has section header but occupies no space

ELF header
Segment header table (required for executables)
.text section
.rodata section
.data section
.bss section
.symtab section
.rel.text section
.rel.data section
.debug section
Section header table

0

ELF Object File Format (cont.)

- .symtab section
 - Symbol table
 - Procedure and static variable names
 - Section names and locations
- .rel.text section
 - Relocation info for .text section
 - Addresses of instructions that will need to be modified in the executable
 - Instructions for modifying.
- .rel.data section
 - Relocation info for .data section
 - Addresses of pointer data that will need to be modified in the merged executable
- .debug section
 - Info for symbolic debugging (gcc -g)
- Section header table
 - Offsets and sizes of each section

ELF header
Segment header table (required for executables)
.text section
.rodata section
.data section
.bss section
.symtab section
.rel.text section
.rel.data section
.debug section
Section header table

Example C Program

m.c

```
int e=7;

int main() {
    int r = a();
    exit(0);
}
```

a.c

```
extern int e;

int *ep=&e;
int x=15;
int y;

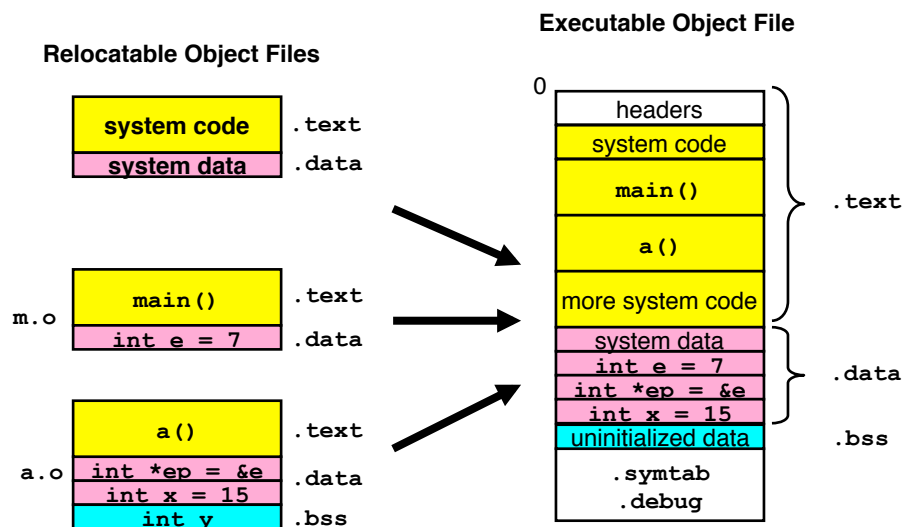
int a() {
    return *ep+x+y;
}
```

5/8/15

CSCI 2021

13

Merging Relocatable Object Files into an Executable Object File



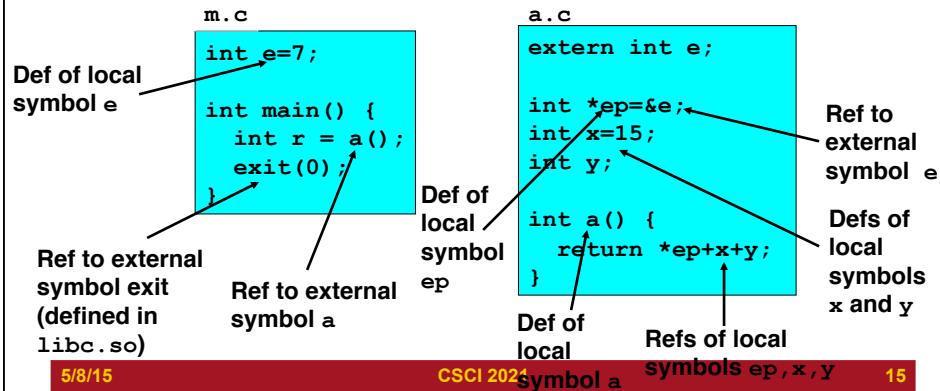
5/8/15

CSCI 2021

14

Relocating Symbols and Resolving External References

- *Symbols* are lexical entities that name functions and variables.
- Each symbol has a *value* (typically a memory address).
- Code consists of symbol *definitions* and *references*.
- References can be either *local* or *external*.



m.o Relocation Info

m.c

```
int e=7;

int main() {
    int r = a();
    exit(0);
}
```

objdump -d
objdump -D
objdump -r

0:	55	push %ebp
1:	89 e5	mov %esp,%ebp
...		
1a:	29 c4	sub %eax,%esp
1c:	e8 fc ff ff	call 1d <main+0x1d>
1d:		R_386_PC32 a
21:	89 45 fc	mov %eax,0xffffffff(%ebp)
24:	c7 04 24 00 00 00	movl \$0x0,(%esp)
2b:	e8 fc ff ff	call 2c <main+0x2c>
2c:		R_386_PC32 exit

Disassembly of section .data:

```
00000000 <e>:
0: 07      pop %es
1: 00 00    add %al,(%eax)
```

RELOCATION RECORDS FOR [.text]:

OFFSET	TYPE	VALUE
0000001d	R_386_PC32	a
0000002c	R_386_PC32	exit

Footer: 5/8/15 CSCI 2021 16

a.o Relocation Info (.text)

a.c

```
extern int e;

int *ep=&e;
int x=15;
int y;

int a() {
    return *ep+x+y;
}
```

Disassembly of section .text:

```
00000000 <a>:
0: 55          push  %ebp
1: 89 e5       mov   %esp,%ebp
3: a1 00 00 00 mov   0x0,%eax
4:           R_386_32  ep
8: 8b 10       mov   (%eax),%edx
a: a1 00 00 00 mov   0x0,%eax
b:           R_386_32  x
f: 01 c2      add   %eax,%edx
11: a1 00 00 00 mov   0x0,%eax
12:           R_386_32  y
16: 8d 04 02    lea   (%edx,%eax,1),%eax
19: 5d         pop   %ebp
1a: c3         ret
```

5/8/15

CSCI 2021

17

a.o Relocation Info (.data)

a.c

```
extern int e;

int *ep=&e;
int x=15;
int y;

int a() {
    return *ep+x+y;
}
```

```
00000000 <ep>:
0: 00 00      add   %al,(%eax)
0:           R_386_32  e
00000004 <x>:
4: 0f 00 00   sldtl (%eax)
...
```

RELOCATION RECORDS FOR [.text]:

OFFSET	TYPE	VALUE
00000004	R_386_32	ep
0000000b	R_386_32	x
00000012	R_386_32	y

RELOCATION RECORDS FOR [.data]:

OFFSET	TYPE	VALUE
00000000	R_386_32	e

5/8/15

CSCI 2021

18

Executable After Relocation and External Reference Resolution

```

08048354 <main>:
...
8048370:  e8 0f 00 00 00      call  8048384 <a>
8048375:  89 45 fc            mov   %eax,0xffffffff(%ebp)
8048378:  c7 04 24 00 00 00 00 movl  $0x0,(%esp)
804837f:  e8 1c ff ff ff      call  80482a0 <exit@plt>

gcc m.o a.o
objdump -d
5/8/15
08048384 <a>:
...
8048387:  a1 ac 95 04 08      mov   0x80495ac,%eax
804838c:  8b 10              mov   (%eax),%edx
804838e:  a1 b0 95 04 08      mov   0x80495b0,%eax
8048393:  01 c2              add   %eax,%edx
8048395:  a1 b8 95 04 08      mov   0x80495b8,%eax
804839a:  8d 04 02           lea   (%edx,%eax,1),%eax
...

```

Executable After Relocation and External Reference Resolution

```

m.c
int e=7;

int main() {
    int r = a();
    exit(0);
}

a.c
extern int e;

int *ep=&e;
int x=15;
int y;

int a() {
    return *ep+x+y;
}

Disassembly of section .data:

080495a8 <c>:
80495a8:  07                pop   %es
80495a9:  00 00            add   %al,(%eax)
...

080495ac <ep>:
80495ac:  a8 95            test  $0x95,%al
80495ae:  04 08            add   $0x8,%al

080495b0 <x>:
80495b0:  0f 00 00          sldtl (%eax)
...

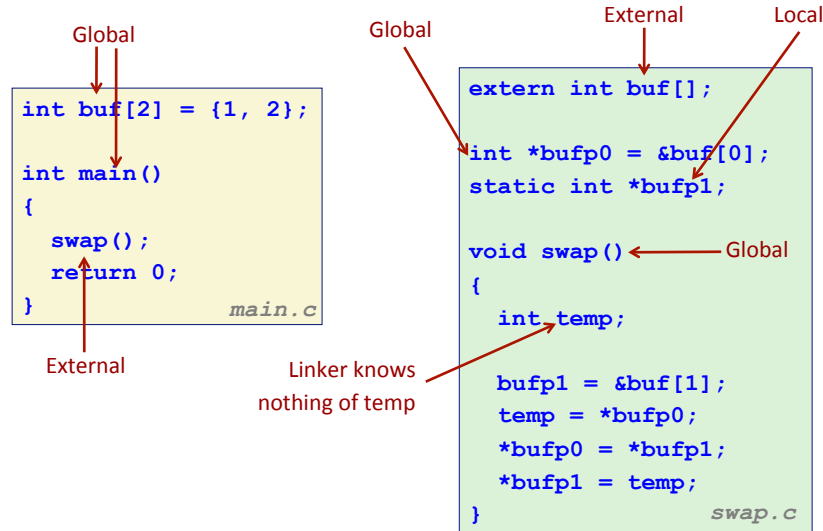
```

5/8/15

CSCI 2021

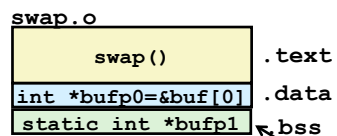
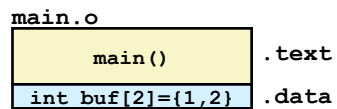
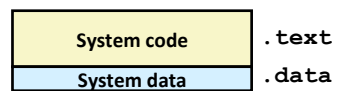
20

Resolving Symbols

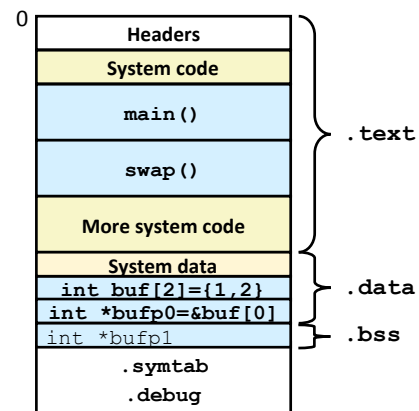


Relocating Code and Data

Relocatable Object Files



Executable Object File



Even though private to swap, requires allocation in .bss

Relocation Info (main)

main.c

```
int buf[2] =
{1,2};

int main()
{
    swap();
    return 0;
}
```

main.o

```
00000000 <main>:
 0: 55                push    %ebp
 1: 89 e5             mov     %esp,%ebp
 3: 83 e4 f0          and     $0xffffffff0,%esp
 6: e8 fc ff ff ff    call    7 <main+0x7>
                    7: R_386_PC32    swap
 b: b8 00 00 00 00    mov     $0x0,%eax
10: 89 ec             mov     %ebp,%esp
12: 5d                pop     %ebp
13: c3                ret
```

```
00000000 <buf>:
 0: 01 00             add     %eax,(%eax)
 2: 00 00             add     %al,(%eax)
 4: 02 00             add     (%eax),%al
 ...
```

Source: objdump -r -d

Relocation Info (swap, .text)

swap.c

```
extern int buf[];

int
*bufp0 = &buf[0];

static int *bufp1;

void swap()
{
    int temp;

    bufp1 = &buf[1];
    temp = *bufp0;
    *bufp0 = *bufp1;
    *bufp1 = temp;
}
```

swap.o

```
Disassembly of section .text:
00000000 <swap>:
 0: 55                push    %ebp
 1: 89 e5             mov     %esp,%ebp
 3: 83 ec 10          sub     $0x10,%esp
 6: c7 05 00 00 00 04 movl    $0x4,0x0
 d: 00 00 00

                    8: R_386_32      .bss
                    c: R_386_32      buf
10: a1 00 00 00 00    mov     0x0,%eax
                    11: R_386_32      bufp0
15: 8b 00             mov     (%eax),%eax
17: 89 45 fc          mov     %eax,-0x4(%ebp)
1a: a1 00 00 00 00    mov     0x0,%eax
                    1b: R_386_32      bufp0
1f: 8b 15 00 00 00 00 mov     0x0,%edx
                    21: R_386_32      .bss
25: 8b 12             mov     (%edx),%edx
27: 89 10             mov     %edx,(%eax)
29: a1 00 00 00 00    mov     0x0,%eax
                    2a: R_386_32      .bss
2e: 8b 55 fc          mov     -0x4(%ebp),%edx
31: 89 10             mov     %edx,(%eax)
33: c9                leave
34: c3                ret
```

Relocation Info (swap, .data)

swap.c

```
extern int buf[];

int *bufp0 = &buf[0];
static int *bufp1;

void swap()
{
    int temp;

    bufp1 = &buf[1];
    temp = *bufp0;
    *bufp0 = *bufp1;
    *bufp1 = temp;
}
```

Disassembly of section .data:

```
00000000 <bufp0>:
  0:  00 00          add    %al, (%eax)
  ...
```

Disassembly of section .bss:

```
00000000 <bufp1>:
  0:  00 00          add    %al, (%eax)
```

Executable Before/After Relocation (.text)

```
08048394 <main>:
8048394:  55              push   %ebp
8048395:  89 e5           mov    %esp, %ebp
8048397:  83 e4 f0        and    $0xfffffff0, %esp
804839a:  e8 09 00 00 00  call   80483a8 <swap>
804839f:  b8 00 00 00 00  mov    $0x0, %eax
80483a4:  89 ec           mov    %ebp, %esp
80483a6:  5d             pop    %ebp
80483a7:  c3             ret
```

Executable Before/After Relocation (.text)

```
80483a8: 55          push    %ebp
80483a9: 89 e5       mov     %esp,%ebp
80483ab: 83 ec 10    sub     $0x10,%esp
80483ae: c7 05 24 a0 04 08 14 movl    $0x804a014,0x804a024
80483b5: a0 04 08
80483b8: a1 18 a0 04 08 mov     0x804a018,%eax
80483bd: 8b 00       mov     (%eax),%eax
80483bf: 89 45 fc    mov     %eax,-0x4(%ebp)
80483c2: a1 18 a0 04 08 mov     0x804a018,%eax
80483c7: 8b 15 24 a0 04 08 mov     0x804a024,%edx
80483cd: 8b 12       mov     (%edx),%edx
80483cf: 89 10       mov     %edx,(%eax)
80483d1: a1 24 a0 04 08 mov     0x804a024,%eax
80483d6: 8b 55 fc    mov     -0x4(%ebp),%edx
80483d9: 89 10       mov     %edx,(%eax)
80483db: c9          leave
```

Executable After Relocation (.data)

Disassembly of section .data:

0804a010 <buf>:

```
804a010: 01 00      add     %eax,(%eax)
804a012: 00 00      add     %al,(%eax)
804a014: 02 00      add     (%eax),%al
...
```

0804a018 <bufp0>:

```
804a018: 10         .byte 0x10
804a019: a0         .byte 0xa0
804a01a: 04 08      add     $0x8,%al
```

Disassembly of section .bss:

...

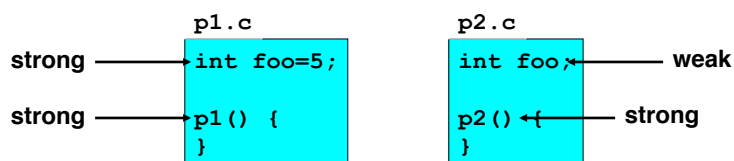
0804a024 <bufp1>:

```
804a024: 00 00      add     %al,(%eax)
```

Strong and Weak Symbols

Program symbols are either strong or weak

- *strong*: procedures and initialized globals
- *weak*: uninitialized globals



5/8/15

CSCI 2021

29

Linker's Symbol Rules

- Rule 1. A strong symbol can only appear once.
- Rule 2. A weak symbol can be overridden by a strong symbol of the same name.
 - references to the weak symbol resolve to the strong symbol.
- Rule 3. If there are multiple weak symbols, the linker can pick an arbitrary one.

5/8/15

CSCI 2021

30

Linker Puzzles

```
int x;
p1() {}
```

```
p1() {}
```

Link time error: two strong symbols (p1)

```
int x;
p1() {}
```

```
int x;
p2() {}
```

References to `x` will refer to the same uninitialized int. Is this what you really want?

```
int x;
int y;
p1() {}
```

```
double x;
p2() {}
```

Writes to `x` in `p2` might overwrite `y`!
Evil!

```
int x=7;
int y=5;
p1() {}
```

```
double x;
p2() {}
```

Writes to `x` in `p2` will overwrite `y`!
Nasty!

```
int x=7;
p1() {}
```

```
int x;
p2() {}
```

References to `x` will refer to the same initialized variable.

Nightmare scenario: two identical weak structs, compiled by different compilers with different alignment rules.

Role of .h Files

c1.c

```
#include "global.h"

int f() {
    return g+1;
}
```

c2.c

```
#include <stdio.h>
#include "global.h"

int main() {
    if (!init)
        g = 37;
    int t = f();
    printf("Calling f yields %d\n", t);
    return 0;
}
```

global.h

```
#ifndef INITIALIZE
int g = 23;
static int init = 1;
#else
int g;
static int init = 0;
#endif
```


Running Preprocessor

c1.c

```
#include "global.h"
int f() {
    return g+1;
}
```

global.h

```
#ifdef INITIALIZE
int g = 23;
static int init = 1;
#else
int g;
static int init = 0;
#endif
```

-DINITIALIZE

no initialization

```
int g = 23;
static int init = 1;
int f() {
    return g+1;
}
```

```
int g;
static int init = 0;
int f() {
    return g+1;
}
```

#include causes C preprocessor to insert file verbatim

Role of .h Files (cont.)

c1.c

```
#include "global.h"

int f() {
    return g+1;
}
```

global.h

```
#ifdef INITIALIZE
int g = 23;
static int init = 1;
#else
int g;
static int init = 0;
#endif
```

c2.c

```
#include <stdio.h>
#include "global.h"

int main() {
    if (!init)
        g = 37;
    int t = f();
    printf("Calling f yields %d\n", t);
    return 0;
}
```

What happens:

gcc -o p c1.c c2.c

??

gcc -o p c1.c c2.c \

-DINITIALIZE

??

Global Variables

- Avoid if you can
- Otherwise
 - Use `static` if you can
 - Initialize if you define a global variable
 - Use `extern` if you use external global variable

Packaging Commonly Used Functions

How to package functions commonly used by programmers?

- Math, I/O, memory management, string manipulation, etc.

Awkward, given the linker framework so far:

- Option 1: Put all functions in a single source file
 - Programmers link big object file into their programs
 - Space and time inefficient
- Option 2: Put each function in a separate source file
 - Programmers explicitly link appropriate binaries into their programs
 - More efficient, but burdensome on the programmer

Packaging Commonly Used Functions

Solution: *static libraries* (.a archive files)

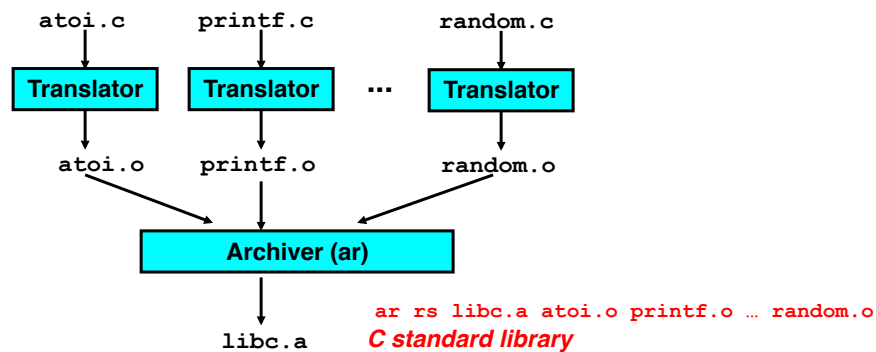
- Concatenate related relocatable object files
 - Concatenate into a single file with an index
 - A.k.a. archive
- Enhance linker
 - Linker tries to resolve unresolved external references
 - Linker looks for the symbols in one or more archives.
- If an archive member file resolves reference,
 - link into executable
- Further improves modularity and efficiency:
 - packaging commonly used functions [e.g., C standard library (libc), math library (libm)]

5/8/15

CSCI 2021

37

Creating Static Libraries



Archiver allows incremental updates:
Recompile function that changes and replace .o file in archive.

5/8/15

CSCI 2021

38

Commonly Used Libraries

`libc.a` (the C standard library)

- 8 MB archive of 900 object files.
- I/O, memory allocation, signal handling, string handling, data and time, random numbers, integer math

`libm.a` (the C math library)

- 1 MB archive of 226 object files.
- floating point math (sin, cos, tan, log, exp, sqrt, ...)

```
% ar -t /usr/lib/libc.a | sort
...
fork.o
...
fprintf.o
fpu_control.o
fputc.o
freopen.o
fscanf.o
fseek.o
fstab.o
...
```

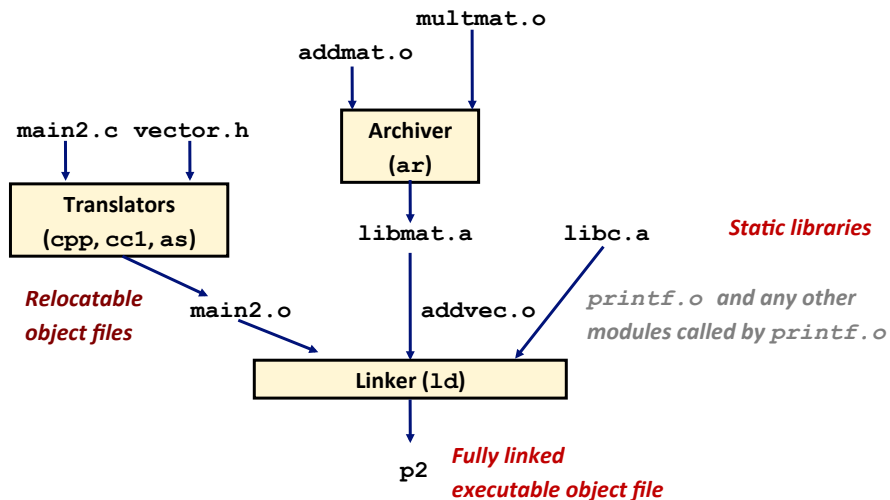
```
% ar -t /usr/lib/libm.a | sort
...
e_acos.o
e_acosf.o
e_acosh.o
e_acoshf.o
e_acoshl.o
e_acosl.o
e_asin.o
e_asinf.o
e_asinl.o
...
```

9/8/19

CSCI 2021

39

Linking with Static Libraries



Using Static Libraries

Linker's algorithm for resolving external references:

- Scan .o files and .a files in the command line order.
- During the scan, keep a list of the current unresolved references.
- As each new .o or .a file obj is encountered, try to resolve each unresolved reference in the list against the symbols in obj.
- If any entries in the unresolved list at end of scan, then error.

Problem:

- Command line order matters!
- Moral: put libraries at the end of the command line.

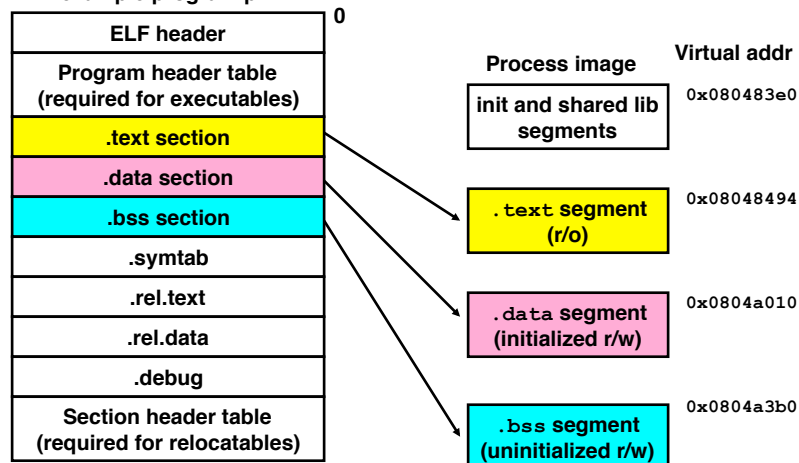
```
lind40-14:/tmp> gcc -m32 -c main.c
lind40-14:/tmp> gcc -m32 -c swap.c
lind40-14:/tmp> ar -q libswap.a swap.o
lind40-14:/tmp> gcc -L. -m32 main.o -lswap
lind40-14:/tmp> gcc -L. -m32 -lswap main.o
main.o: In function `main':
main.c:(.text+0x7): undefined reference to `swap'
collect2: ld returned 1 exit status
```

5/8/15

41

Loading Executable Binaries

Executable object file for
example program p

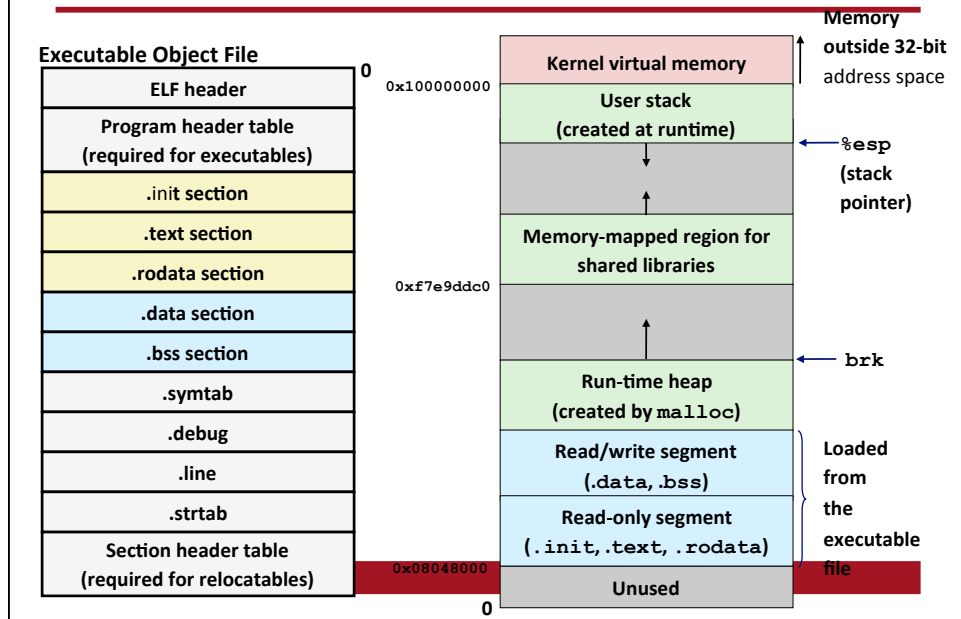


5/8/15

CSCI 2021

42

Loading Executable Object Files



Shared Libraries

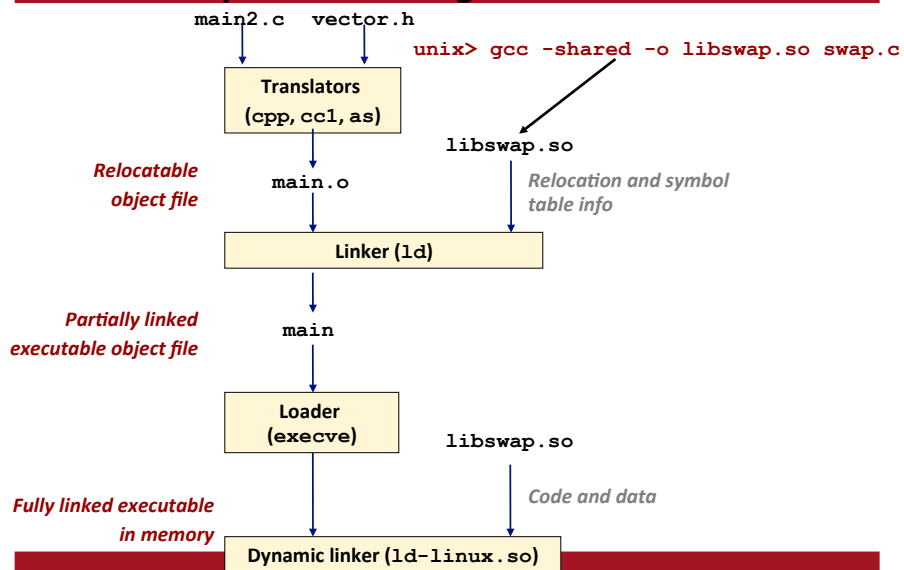
Static libraries have the following disadvantages:

- Potential for duplication
 - lots of common code in the executable files on a filesystem
 - e.g., every C program needs the standard C library
- Potential for duplication
 - lots of code in the virtual memory space of many processes.
- Relink
 - Minor bug fixes of system libraries require each application to explicitly relink

Shared Libraries (cont.)

- Modern solution: Shared Libraries
 - Object files that contain code and data that are loaded and linked into an application *dynamically*, at either *load-time* or *run-time*
 - Also called: dynamic link libraries, DLLs, `.so` files
- Dynamic linking can occur when executable is first loaded and run (load-time linking).
 - Common case for Linux, handled automatically by the dynamic linker (`ld-linux.so`).
 - Standard C library (`libc.so`) usually dynamically linked.
- Dynamic linking can also occur after program has begun (run-time linking).
 - In Linux, this is done by calls to the `dlopen()` interface.
 - Distributing software.
 - High-performance web servers.
 - Runtime library interpositioning.
- Shared library routines can be shared by multiple processes.
 - More on this when we learn about virtual memory

Dynamic Linking at Load-time



Dynamic Linking at Run-time

```
#include <stdio.h>
#include <dlfcn.h>

int x[2] = {1, 2};
int y[2] = {3, 4};
int z[2];

int main() {
    void *handle;
    void (*addvec)(int *, int *, int *, int);
    char *error;

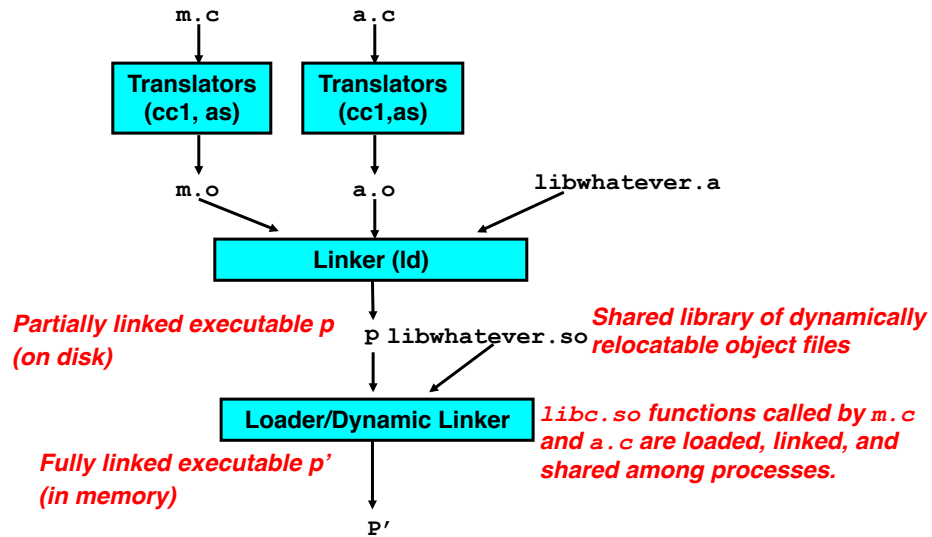
    /* dynamically load the shared lib that contains addvec() */
    handle = dlopen("./libvector.so", RTLD_LAZY);
    if (!handle) {
        fprintf(stderr, "%s\n", dlerror());
        exit(1);
    }
}
```

Dynamic Linking at Run-time

```
...
/* get a pointer to the addvec() function we just loaded */
addvec = dlsym(handle, "addvec");
if ((error = dlerror()) != NULL) {
    fprintf(stderr, "%s\n", error);
    exit(1);
}
/* Now we can call addvec() just like any other function */
addvec(x, y, z, 2);
printf("z = [%d %d]\n", z[0], z[1]);

/* unload the shared library */
if (dlclose(handle) < 0) {
    fprintf(stderr, "%s\n", dlerror());
    exit(1);
}
return 0;
}
```


The Complete Picture



5/8/15

CSCI 2021

49