


Database Management Systems (DBMS)

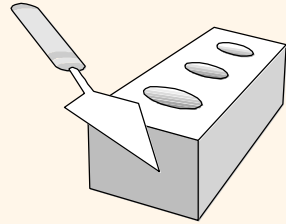
Chapter 1

Instructor:

Mohamed Mokbel (mokbel@cs.umn.edu)



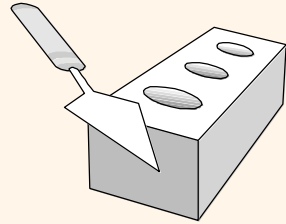
Welcome to CSCI 4707



❖ Instructor:

- Mohamed F. Mokbel
Office: KHKH 4-207
Web: www.cs.umn.edu/~mokbel
Email: mokbel@cs.umn.edu **(Include [CSCI 4707] on the Subject line)**
Office Hours: Tu: 12:00 PM – 1:00 PM
 Thu: 2:15 PM – 3:15 PM
 (or by appointment)


Welcome to CSCI 4707



❖ TA:


- Christopher Jonathan
Office KHKH: 2-209
Email: cjonathan@umn.edu
Office hours: Tue : 4:00 PM - 6:00 PM
- Ibrahim Sabek
Office KHKH: 2-209
Email: sabek@cs.umn.edu
Office hours: Wed : 11:00 AM - 1:00 PM

Grading Policy



- ❖ Final Exam: 30% (Inclusive)
- ❖ Two Midterm Exams: 25% (12.5% each)
- ❖ Four Homeworks: 10% (2.5% each)
- ❖ Four Labs: 35%
 - Lab # 1 (3%): Individual
 - Lab # 2 (7%): Individual
 - Lab # 3 (13%): (Group of 2)
 - Lab # 4 (12%): (Group of 2)
- ❖ Every one-day late submission reduces 10% of the grade.
- ❖ No more than three-day late submission is allowed for each lab/homework, no more than 6 late submission days for all labs/homeowrks.

Welcome to CSCI 4707



❖ Text Book:

Database Management Systems,
3e (ISBN: 0-07-246563-8)

Raghu Ramakrishnan and Johannes
Gehrke


Book website:

<http://www.cs.wisc.edu/~dbbook/>


- Chapters: 1-5, 8, 9, 12-20

❖ Class web page:

- <http://www-users.cselabs.umn.edu/classes/Fall-2015/csci4707/index.php>
- Tentative scheduling is in the course web site
- Stay tuned for announcements, labs, homeworks, and grades




What we will study in the Course



- ❖ **PART 1:** *Outside* the Database Engine as a *User*
 - How to do conceptual database design (Chapter 2)
 - How to do a logical database design (Chapter 3)
 - How to query the database (Chapters 4, 5)
- ❖ **PART 2:** *Inside* the Database Engine
 - Storage and indexing modules (Chapters 8-10)
 - Query processing & optimization (Chapters 12-15)
 - Transaction Processing (Chapters 16-18)
- ❖ **PART 3:** *Outside* the Database Engine as a *DBA*
 - How to refine your schema design (Chapter 19)
 - How to tune the database design (Chapter 20)

What Is a DBMS?




- ❖ How do you store your friend's phone numbers:

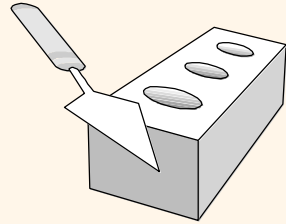
Name	Home	Cell	Address	Email

- ❖ Designing, storing, and managing such “simple” tables is the core of DBMS

What Is a DBMS?



- ❖ A **VERY LARGE**, integrated collection of data.
- ❖ Models real-world:
 - Entities (e.g., universities, departments, students, courses)
 - Relationships (e.g., Students are taking courses)
- ❖ A Database Management System (DBMS) is a software package designed to store and manage data.



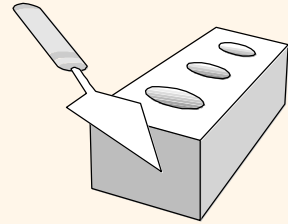
History of DBMS

❖ Early 60's

- First general-purpose DBMS by Charles Bachman at General Electric (First recipient of **ACM Turing Award** in 1973)

❖ Late 60's

- Hierarchical data model developed at IBM
- The SABRE system for airline reservation is jointly developed by American Airlines and IBM where several people can access the same data through a network



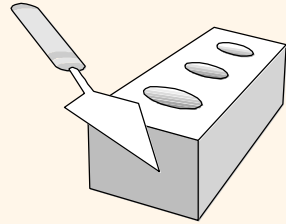
History of DBMS

❖ 70's

- The relational data model is proposed by Edgar Codd at IBM (**ACM Turing Award** at 1981)
- Two main prototypes for relational database management systems are developed. Ingres at UCB and System R at IBM (Mike Stonebraker received the **ACM Turing Award** at 2014)
- Peter Chen (MIT) proposed the entity-relationship model

❖ 80's

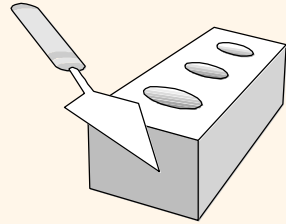
- SQL query language is developed (part of System R)
- The concept of read/write transactions is developed to allow concurrent execution of database operations (Jim Gray received the **ACM Turing Award** at 1998)
- Commercial databases are in the market (DB2, Oracle, Informix)



History of DBMS

❖ 90's - present


- DBMSs are well-established in industry and academia
- New database applications:
 - Data warehousing
 - Multimedia databases
 - Spatial/Spatio-temporal databases
 - Data mining
 - Scientific database
 - Bioinformatics
 - Digital libraries
 - Web-databases



Files vs. DBMS


- ❖ Special code for different queries
- ❖ Must protect data from inconsistency due to multiple concurrent users
- ❖ Crash recovery
- ❖ Security and access control

Why Use a DBMS?




- ❖ **Data independence**
 - Applications are independent from data representation
- ❖ **Efficient data access**
 - Through indexing and query optimization techniques
- ❖ **Data integrity and security**
 - DBMS enforce integrity constraints and access control
- ❖ **Concurrent access**
 - Multiples users are allowed to use the same tables
- ❖ **Crash recovery**
 - DBMS protects the user from the system failure
- ❖ **Reduced application development time**
 - DBMS supports functions that are common to many applications

Why Study Databases??




- ❖ Shift from computation to information
 - at the “low end”: scramble to webspace (a mess!)
 - at the “high end”: scientific applications
- ❖ Datasets increasing in diversity and volume.
 - Digital libraries, interactive video, Human Genome project, Earth Observation project
 - ... need for DBMS exploding
- ❖ Good job market as a DBA, system analyst,
- ❖ 3-credit course(s)



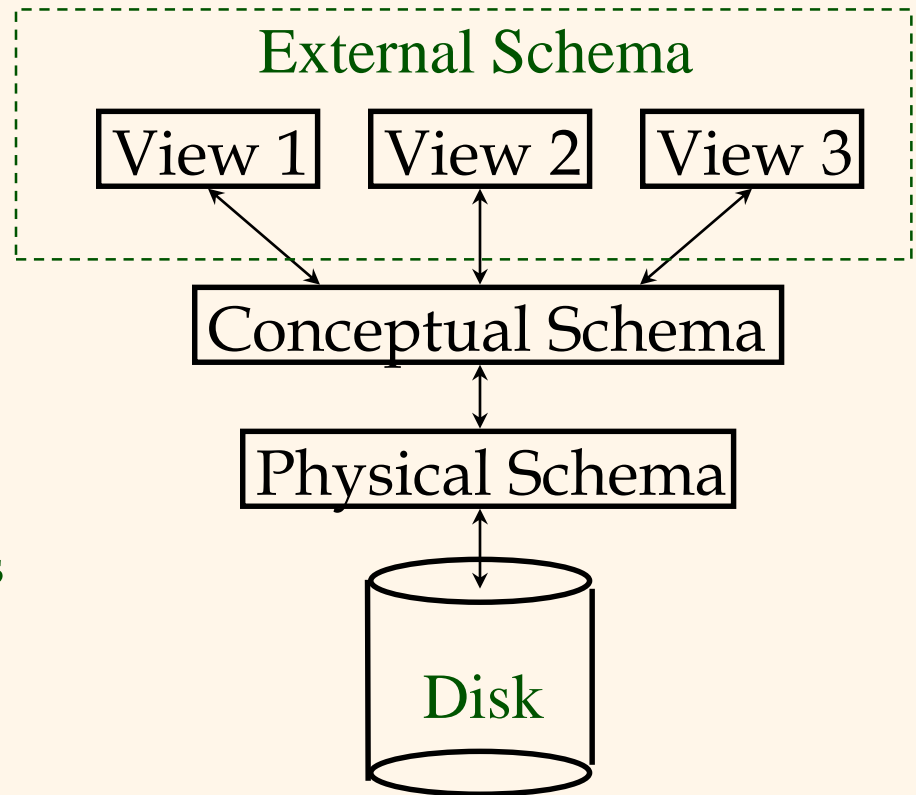
Concepts of Data Modeling

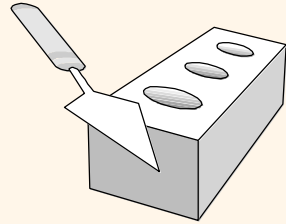
- ❖ A data model is a collection of concepts for describing data.
- ❖ A schema is a description of a particular collection of data, using the a given data model.
- ❖ The relational model of data is the most widely used model today.
 - Main concept: relation, basically a table with rows and columns.
 - Every relation has a schema, which describes the columns, or fields.



Levels of Abstraction

- ❖ Many views, single conceptual schema and physical schema.
 - Views (**External schema**) describe how users see the data.
 - **Conceptual schema** defines logical structure
 - **Physical schema** describes the files and indexes used.





Example: University Database

❖ Conceptual schema:

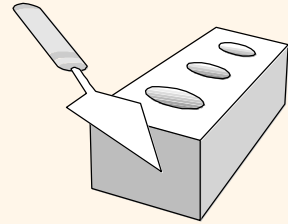
- *Students(sid: string, name: string, login: string, age: integer, gpa: real)*
- *Courses(cid: string, cname: string, credits: integer)*
- *Enrolled(sid: string, cid: string, grade: string)*

❖ Physical schema:

- Relations stored as unordered files.
- Index on first column of Students.

❖ External Schema (View):


- *Course_info(cid: string, enrollment: integer)*



*Data Independence **


- ❖ Applications insulated from how data is structured and stored.
- ❖ Logical data independence: Protection from changes in *logical* structure of data.
- ❖ Physical data independence: Protection from changes in *physical* structure of data.

** One of the most important benefits of using a DBMS!*



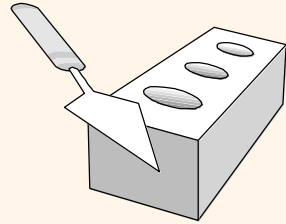
Concurrency Control

- ❖ Concurrent execution of user programs is essential for good DBMS performance.
 - Because disk accesses are frequent, and relatively slow, it is important to keep the CPU working on several user programs concurrently.
- ❖ Interleaving actions of different user programs can lead to inconsistency: e.g., check is cleared while account balance is being computed.
- ❖ DBMS ensures such problems don't arise: users can pretend they are using a single-user system.



Transaction: An Execution of a DB Program

- ❖ Key concept is transaction, which is an *atomic* sequence of database actions (reads/writes).
- ❖ Each transaction, executed completely, must leave the DB in a consistent state if DB is consistent when the transaction begins.
 - Users can specify some simple integrity constraints on the data, and the DBMS will enforce these constraints.
 - Beyond this, the DBMS does not really understand the semantics of the data. (e.g., it does not understand how the interest on a bank account is computed).
 - Thus, ensuring that a transaction (run alone) preserves consistency is ultimately the *user's* responsibility!

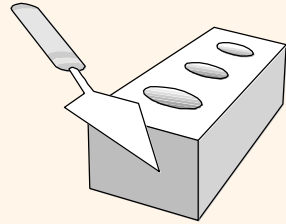


Example

- ❖ Consider two transactions (*Xacts*):

T1:	BEGIN	$A=A+100$,	$B=B-100$	END
T2:	BEGIN	$A=1.06*A$,	$B=1.06*B$	END

- ❖ Intuitively, the first transaction is transferring \$100 from B's account to A's account. The second is crediting both accounts with a 6% interest payment.
- ❖ There is no guarantee that T1 will execute before T2 or vice-versa, if both are submitted together. However, the net effect *must* be equivalent to these two transactions running serially in some order.



Example (Contd.)

- ❖ Consider a possible interleaving (schedule):

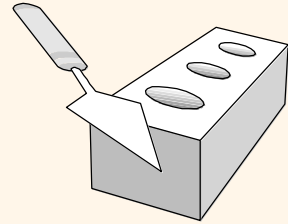
T1:	$A = A + 100,$	$B = B - 100$
T2:	$A = 1.06 * A,$	$B = 1.06 * B$

- ❖ This is OK. But what about:

T1:	$A = A + 100,$	$B = B - 100$
T2:	$A = 1.06 * A, B = 1.06 * B$	

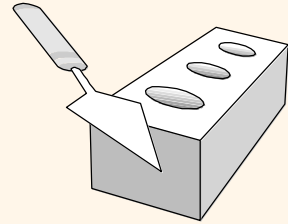
- ❖ The DBMS's view of the second schedule:

T1:	$R(A), W(A),$	$R(B), W(B)$
T2:	$R(A), W(A), R(B), W(B)$	



Scheduling Concurrent Transactions

- ❖ DBMS ensures that execution of $\{T_1, \dots, T_n\}$ is equivalent to some serial execution $T_1' \dots T_n'$.
 - Before reading/writing an object, a transaction requests a lock on the object, and waits till the DBMS gives it the lock. All locks are released at the end of the transaction. (Strict 2PL locking protocol.)
 - **Idea:** If an action of T_i (say, writing X) affects T_j (which perhaps reads X), one of them, say T_i , will obtain the lock on X first and T_j is forced to wait until T_i completes; this effectively orders the transactions.
 - What if T_j already has a lock on Y and T_i later requests a lock on Y ? (Deadlock!) T_i or T_j is aborted and restarted!

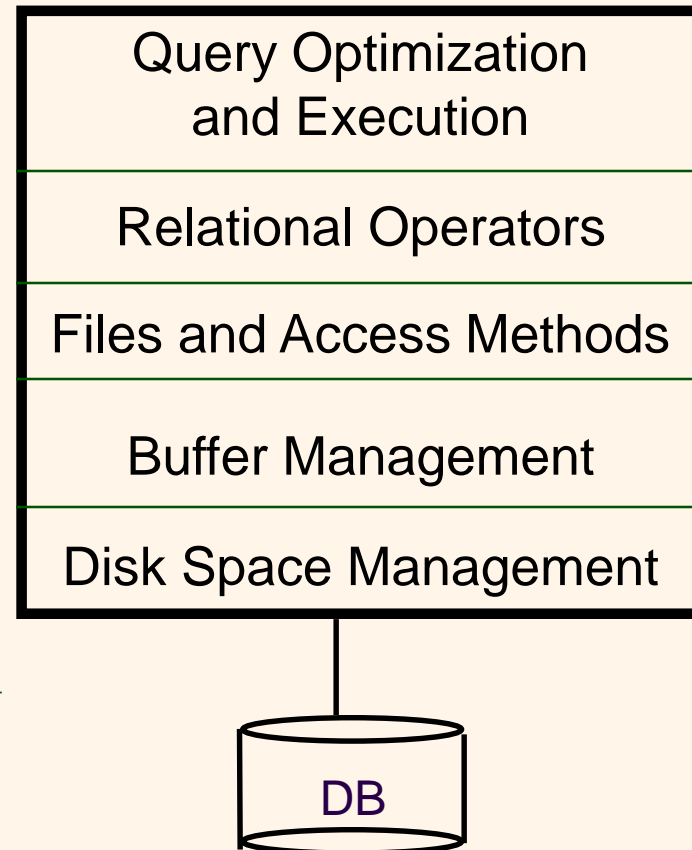


Ensuring Atomicity

- ❖ DBMS ensures *atomicity* (all-or-nothing property) even if system crashes in the middle of a transaction.
- ❖ **Idea:** Keep a log (history) of all actions carried out by the DBMS while executing a set of transactions:
 - **Before** a change is made to the database, the corresponding log entry is forced to a safe location.
 - After a crash, the effects of partially executed transactions are undone using the log.


Structure of a DBMS

- ❖ A typical DBMS has a layered architecture.
- ❖ The figure does not show the concurrency control and recovery components.
- ❖ This is one of several possible architectures; each system has its own variations.




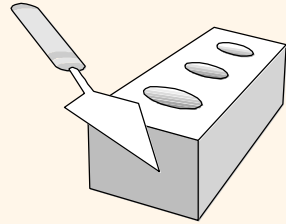
These layers must consider concurrency control and recovery

Databases make these folks happy ...



- ❖ End users
- ❖ DBMS vendors
- ❖ DB application programmers
- ❖ Database administrator (DBA)
 - Designs logical / physical schemas
 - Handles security and authorization
 - Data availability, crash recovery
 - Database tuning as needs evolve





Steps for Designing a Database

1. Requirement analysis

- ❖ An informal discussion with the customers
- ❖ Understand what are the requirements, how different entities relate to each other, what are the frequent operations to be performed


2. Conceptual Database Design

- ❖ Develop a high-level description of the data
- ❖ Develop an ER (entity-relationship) model that capture the semantics of the data

3. Logical Database Design

- ❖ Convert the ER model into a (relational) database schema

Steps for Designing a Database (Cont.)



4. Scheme Refinement


- ❖ Identify potential problems in your schema design
- ❖ This process can be guided by some elegant and powerful theory

5. Physical Database Design

- ❖ Study the expected workload of the system
- ❖ Tune the performance by building indexes and clustering tables

6. Application and Security Design

- ❖ Identify which parts of the database are accessible to whom



Summary

- ❖ DBMS used to maintain, query large datasets.
- ❖ Benefits include recovery from system crashes, concurrent access, quick application development, data integrity and security.
- ❖ Levels of abstraction give data independence.
- ❖ A DBMS typically has a layered architecture.
- ❖ DBAs hold responsible jobs and are **well-paid!**
- ❖ DBMS R&D is one of the broadest, most exciting areas in CS.

