



The Relational Model

Chapter 3



Why Study the Relational Model?


- ❖ Most widely used model.
 - Vendors: IBM, Informix, Microsoft (Access and SQL Server), Oracle, Sybase, FoxBase etc.
- ❖ “Legacy systems” in older models
 - E.G., IBM’s IMS
- ❖ Recent competitor: object-oriented model
 - ObjectStore, Versant, Ontos
 - A synthesis emerging: *object-relational model*
 - Informix Universal Server, UniSQL, O2, Oracle, DB2
 - More in scientific computing



Relational Database: Definitions


- ❖ *Relational database*: a set of *relations*
- ❖ *Relation*: made up of 2 parts:
 - *Relation Schema* : specifies name of relation, plus name and type of each column.
 - E.G. Students(*sid*: string, *name*: string, *login*: string, *age*: integer, *gpa*: real).
 - *Relation Instance* : a *table*, with rows and columns.
#Rows = *cardinality*, #fields = *degree / arity*.
- ❖ Can think of a relation as a *set* of rows or *tuples* or *records* (i.e., all rows are distinct).

Example Instance of Students Relation




sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

❖ Cardinality = 3, degree = 5, all rows distinct



The SQL Query Language


- ❖ Developed by IBM (system R) in the 1970s
- ❖ The most widely used language for **creating**, **manipulating**, and **querying** relational DBMS.
- ❖ Need for a standard since it is used by many vendors
- ❖ Standards:
 - SQL-86
 - SQL-89 (minor revision)
 - SQL-92 (major revision)
 - SQL-99 (major extensions, current standard)



The SQL Query Language

- ❖ A major strength of the relational model: supports simple, powerful *querying* of data.
- ❖ **Queries** can be written intuitively, and the DBMS is responsible for efficient evaluation.
 - The key: precise semantics for relational queries.
 - Allows the optimizer to extensively re-order operations, and still ensure that the answer does not change.

The SQL Query Language (Simple Query)



sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

❖ To find the names and login of all 18 year old students:


```
SELECT name, login  
FROM Students  
WHERE age=18
```

name	login
Jones	jones@cs
Smith	smith@eecs

• To find all student data, replace the first line:

```
SELECT *  
FROM Students  
WHERE age=18
```

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2



Querying Multiple Relations

Given the following instances of Enrolled and Students, list all the students who get grade 'A' in any course along with the course id: :


sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

sid	cid	grade
53831	Carnatic101	C
53831	Reggae203	B
53650	Topology112	A
53666	History105	B

Answer:

S.name	E.cid
Smith	Topology112

```
SELECT S.name, E.cid
FROM Students S, Enrolled E
WHERE S.sid=E.sid AND E.grade="A"
```





Creating Relations in SQL

```
CREATE TABLE Students (sid: CHAR(20),  
                        name: CHAR(20),  
                        login: CHAR(10),  
                        age: INTEGER,  
                        gpa: REAL)
```

- ❖ Creates the Students relation.
 - The type (**domain**) of each field is specified, and enforced by the DBMS whenever tuples are added or modified.

sid	name	login	age	gpa
-----	------	-------	-----	-----




Creating Relations in SQL

```
CREATE TABLE Enrolled (sid: CHAR(20),  
                        cid: CHAR(20),  
                        grade: CHAR(2))
```

- ❖ The Enrolled table holds information about courses that students take.

sid	cid	grade
-----	-----	-------

Destroying and Altering Relations




`DROP TABLE Students`

- ❖ Destroys the relation Students. The schema information *and* the tuples are deleted.

`ALTER TABLE Students`

`ADD COLUMN firstYear: integer`

- ❖ The schema of Students is altered by adding a new field; every tuple in the current instance is extended with a *null* value in the new field.



Adding and Deleting Tuples


- ❖ Can insert a single tuple using:

```
INSERT  
INTO Students (sid, name, login, age, gpa)  
VALUES (53688, 'Smith', 'smith@ee', 18, 3.2)
```

**The list of columns is optional*

- ❖ Can delete all tuples satisfying some condition (e.g., name = Smith):

```
DELETE  
FROM Students S  
WHERE S.name = 'Smith'
```




Updating Tuples

❖ Update a single tuple

```
UPDATE Students S  
SET S.age = S.age +1, S.gpa = S.gpa -1  
WHERE S.sid = 54832
```


❖ Update multiple tuples

```
UPDATE Students S  
SET S.age = S.age +1, S.gpa = S.gpa -1  
WHERE S.gpa > 3.3
```



Integrity Constraints (ICs)


- ❖ **IC:** condition that must be true for *any* instance of the database; e.g., *domain constraints*.
 - ICs are specified when schema is defined.
 - ICs are checked when relations are modified.
- ❖ A *legal* instance of a relation is one that satisfies all specified ICs.
 - DBMS should not allow illegal instances.
- ❖ If the DBMS checks ICs, stored data is more faithful to real-world meaning.
 - Avoids data entry errors, too!



Primary Key Constraints

- ❖ A set of fields is a key for a relation if :
 1. No two distinct tuples can have same values in all key fields, and
 2. This is not true for any subset of the key.
 - Part 2 false? A *superkey*.
 - If there are more than one key for a relation, one of the keys is chosen (by DBA) to be the *primary key* while other keys are considered *candidate key*
- ❖ E.g., *sid* is a key for Students. (What about *name*?) The set {*sid*, *gpa*} is a superkey.

Primary and Candidate Keys in SQL




- ❖ “For a given student and course, there is a single grade.”

```
CREATE TABLE Enrolled (sid CHAR(20)  
                        cid CHAR(20),  
                        grade CHAR(2),  
                        PRIMARY KEY (sid,cid) )
```

- ❖ “Students can take only one course, and receive a single grade for that course; further, no two students in a course receive the same grade.”


```
CREATE TABLE Enrolled (sid CHAR(20)  
                        cid CHAR(20),  
                        grade CHAR(2),  
                        PRIMARY KEY (sid),  
                        UNIQUE (cid, grade) )
```


Primary and Candidate Keys in SQL



- ❖ Possibly many candidate keys (specified using **UNIQUE**), one of which is chosen as the *primary key* (specified using **PRIMARY KEY**) .
- ❖ Used carelessly, an IC can prevent the storage of database instances that arise in practice!

Enforcing Primary Key Constrains




Students


sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

- ❖ Having this table: The following transactions are **REJECTED**:
 - INSERT INTO Students (sid, name, login, age, gpa)
VALUES (53688, 'Mike', 'Mike@cs', 17, 3.4)
 - INSERT INTO Students (sid, name, login, age, gpa)
VALUES (null, 'Mike', 'Mike@cs', 17, 3.4)
 - INSERT INTO Students (sid, name, login, age, gpa)
VALUES (53784, 'Mike', 'Mike@cs', 'twenty', 3.4)
 - UDPATE Students S SET S.sid = 53688
WHERE S.sid = 53666
- ❖ There are no problems for *Deletion*

Foreign Keys, Referential Integrity



- ❖ Foreign key : Set of fields in one relation that is used to `refer' to a tuple in another relation. (Must correspond to primary key of the second relation.) Like a `logical pointer'.
- ❖ E.g. *sid* is a foreign key referring to **Students**:
 - Enrolled(*sid*: string, *cid*: string, *grade*: string)
 - If all foreign key constraints are enforced, referential integrity is achieved, i.e., no dangling references.
- ❖ Can you name a data model w/o referential integrity?



Foreign Keys in SQL

- ❖ Only students listed in the Students relation should be allowed to enroll for courses.


```
CREATE TABLE Enrolled (sid CHAR(20),  
                        cid CHAR(20),  
                        grade CHAR(2),  
                        PRIMARY KEY (sid,cid),  
                        FOREIGN KEY (sid) REFERENCES Students )
```

Enrolled

sid	cid	grade
53666	Carnatic101	C
53666	Reggae203	B
53650	Topology112	A
53666	History105	B

Students

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8



Enforcing Referential Integrity

Enrolled

sid	cid	grade
53666	Carnatic101	C
53666	Reggae203	B
53650	Topology112	A
53666	History105	B

Students


sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

- ❖ What should be done if an Enrolled tuple with a non-existent student id is inserted?

➤ INSERT INTO Enrolled (sid, cid, grade)
VALUES (53611, 'History105', 'A')

- *(Reject it!)*

Enforcing Referential Integrity



Enrolled


sid	cid	grade
53666	Carnatic101	C
53666	Reggae203	B
53650	Topology112	A
53666	History105	B

Students

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

- ❖ What should be done if a Students tuple is deleted?
 - **OPTION 1:** Also delete all Enrolled tuples that refer to it.
 - **OPTION 2:** Disallow deletion of a Students tuple that is referred to.
 - **OPTION 3:** Set sid in Enrolled tuples that refer to it to a *default sid*.
 - **OPTION 4:** Set sid in Enrolled tuples that refer to it to *null*.
- ❖ Similar if primary key of Students tuple is updated.


Referential Integrity in SQL




- ❖ SQL/92 and SQL:1999 support all four options on deletes and updates.
 - Default is **NO ACTION** (*delete/update is rejected*)
 - **CASCADE** (also delete all tuples that refer to deleted tuple)
 - **SET NULL / SET DEFAULT** (sets foreign key value of referencing tuple)

```
CREATE TABLE Enrolled (sid CHAR(20),  
                        cid CHAR(20),  
                        grade CHAR(2),  
                        PRIMARY KEY (sid,cid),  
                        FOREIGN KEY (sid)  
                        REFERENCES Students  
                        ON DELETE CASCADE  
                        ON UPDATE SET DEFAULT )
```

Logical DB Design: ER to Relational





❖ Entity sets to tables:



```
CREATE TABLE Employees (ssn CHAR(11),  
                          name CHAR(20),  
                          lot INTEGER,  
                          PRIMARY KEY (ssn))
```

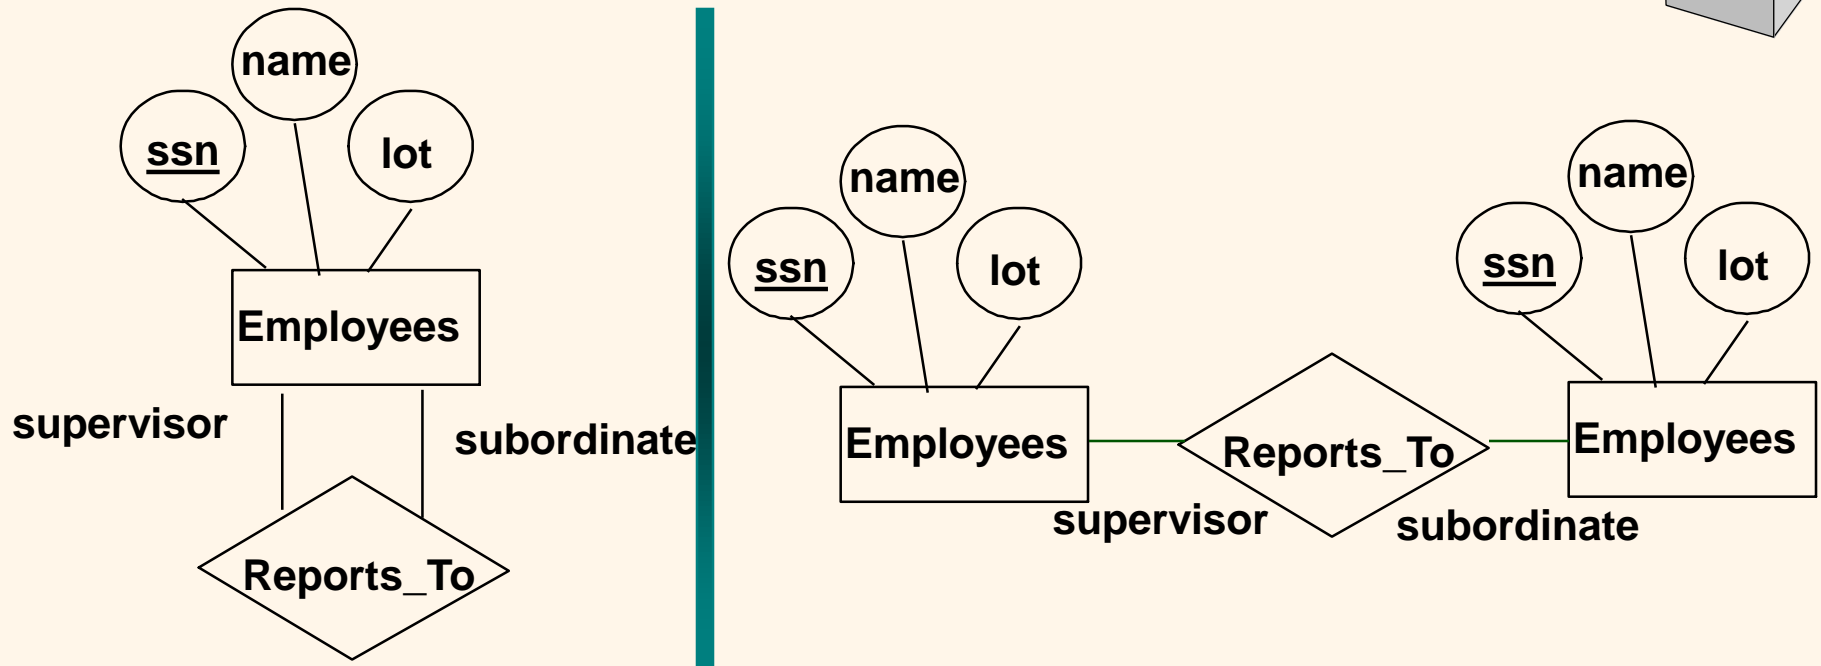


Relationship Sets to Tables



- ❖ In translating a relationship set to a relation, attributes of the relation must include:
 - Keys for each participating entity set (as foreign keys).
 - All descriptive attributes.


```
CREATE TABLE Works_In( ssn CHAR(11),  
                        did INTEGER,  
                        since DATE,  
                        PRIMARY KEY (ssn, did),  
                        FOREIGN KEY (ssn) REFERENCES Employees,  
                        FOREIGN KEY (did) REFERENCES Departments)
```

Relationship Sets to Tables





```
CREATE TABLE Reports_To( supervisor_ssn CHAR(11),
                        subordinate_ssn CHAR(11),
PRIMARY KEY (supervisor_ssn, subordinate_ssn),
FOREIGN KEY (supervisor_ssn) REFERENCES Employees(ssn),
FOREIGN KEY (subordinate_ssn) REFERENCES Employees(ssn))
```


Review: Key Constraints




- ❖ Each dept has at most one manager, according to the key constraint on Manages.




1-to-1



1-to Many




Many-to-1



Many-to-Many

Translation to relational model?



Translating ER Diagrams with Key Constraints


❖ Map relationship to a table:

- Separate tables for Employees and Departments.

```
CREATE TABLE Manages(  
  ssn CHAR(11),  
  did INTEGER,  
  since DATE,  
  PRIMARY KEY (ssn,did),  
  FOREIGN KEY (ssn) REFERENCES Employees,  
  FOREIGN KEY (did) REFERENCES Departments)
```

❖ Note that **did** is the key now!

```
CREATE TABLE Manages(  
  ssn CHAR(11),  
  did INTEGER,  
  since DATE,  
  PRIMARY KEY (did),  
  FOREIGN KEY (ssn) REFERENCES Employees,  
  FOREIGN KEY (did) REFERENCES Departments)
```





Translating ER Diagrams with Key Constraints


- ❖ Since each department has a unique manager, we could instead combine Manages and Departments.

```
CREATE TABLE Dept_Mgr(  
  did INTEGER,  
  dname CHAR(20),  
  budget REAL,  
  ssn CHAR(11),  
  since DATE,  
  PRIMARY KEY (did),  
  FOREIGN KEY (ssn) REFERENCES Employees)
```

Review: Participation Constraints



- ❖ Does every department have a manager?
 - If so, this is a participation constraint: the participation of Departments in Manages is said to be *total* (vs. *partial*).
 - Every *did* value in Departments table must appear in a row of the Manages table (with a non-null *ssn* value!)





Participation Constraints in SQL


- ❖ We can capture participation constraints involving one entity set in a binary relationship, but little else (without resorting to CHECK constraints).

```
CREATE TABLE Dept_Mgr(  
    did INTEGER,  
    dname CHAR(20),  
    budget REAL,  
    ssn CHAR(11) NOT NULL,  
    since DATE,  
    PRIMARY KEY (did),  
    FOREIGN KEY (ssn) REFERENCES Employees,  
    ON DELETE NO ACTION)
```

Review: Weak Entities



- ❖ A *weak entity* can be identified uniquely only by considering the primary key of another (*owner*) entity.
 - Owner entity set and weak entity set must participate in a one-to-many relationship set (1 owner, many weak entities).
 - Weak entity set must have total participation in this *identifying* relationship set.




Translating Weak Entity Sets

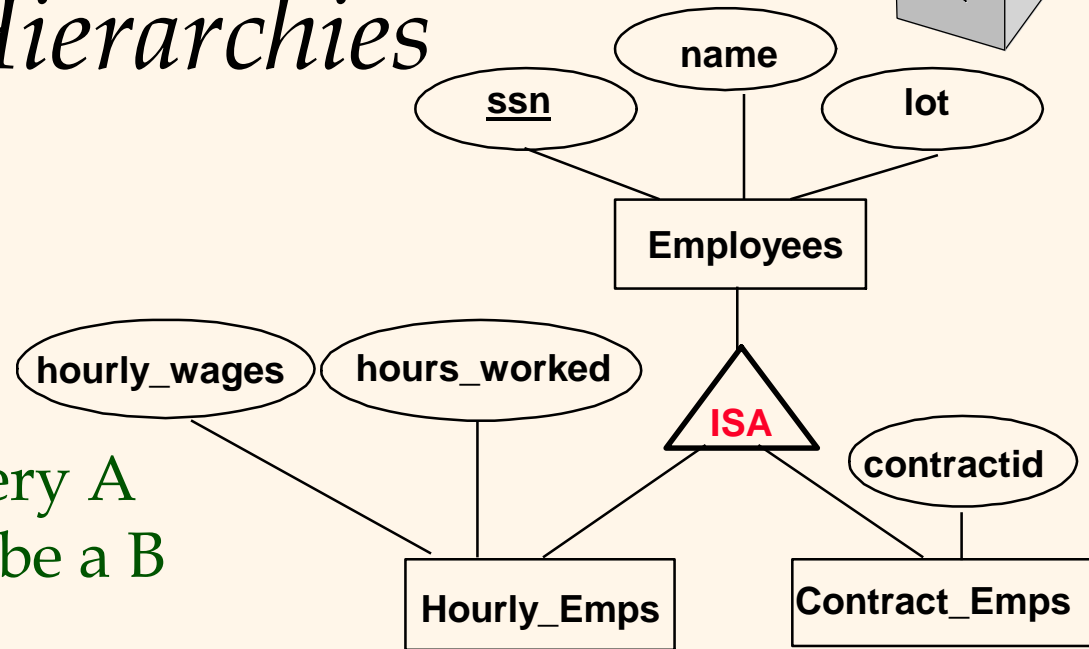
- ❖ Weak entity set and identifying relationship set are translated into a single table.
 - When the owner entity is deleted, all owned weak entities must also be deleted.

```
CREATE TABLE Dep_Policy (  
    pname CHAR(20),  
    age INTEGER,  
    cost REAL,  
    ssn CHAR(11) NOT NULL,  
    PRIMARY KEY (pname, ssn),  
    FOREIGN KEY (ssn) REFERENCES Employees,  
    ON DELETE CASCADE)
```


Review: ISA Hierarchies



- ❖ As in C++, or other PLs, attributes are inherited.
- ❖ If we declare A **ISA** B, every A entity is also considered to be a B entity.



- ❖ **Overlap constraints:** Can Joe be an Hourly_Emps as well as a Contract_Emps entity? (*Allowed/disallowed*)
 - Hourly_Emps OVERLAPS Senior_Emps
- ❖ **Covering constraints:** Does every Employees entity also have to be an Hourly_Emps or a Contract_Emps entity? (*Yes/no*)
 - Hourly_Emps AND Contract_Emps COVERS Employees



Translating ISA Hierarchies to Relations


❖ *General approach:*

- 3 relations: Employees, Hourly_Emps and Contract_Emps.
 - *Hourly_Emps*: Every employee is recorded in Employees. For hourly emps, extra info recorded in Hourly_Emps (*hourly_wages*, *hours_worked*, *ssn*); must delete Hourly_Emps tuple if referenced Employees tuple is deleted).
 - Queries involving all employees easy, those involving just Hourly_Emps require a join to get some attributes.

❖ *Alternative: Just Hourly_Emps and Contract_Emps.*

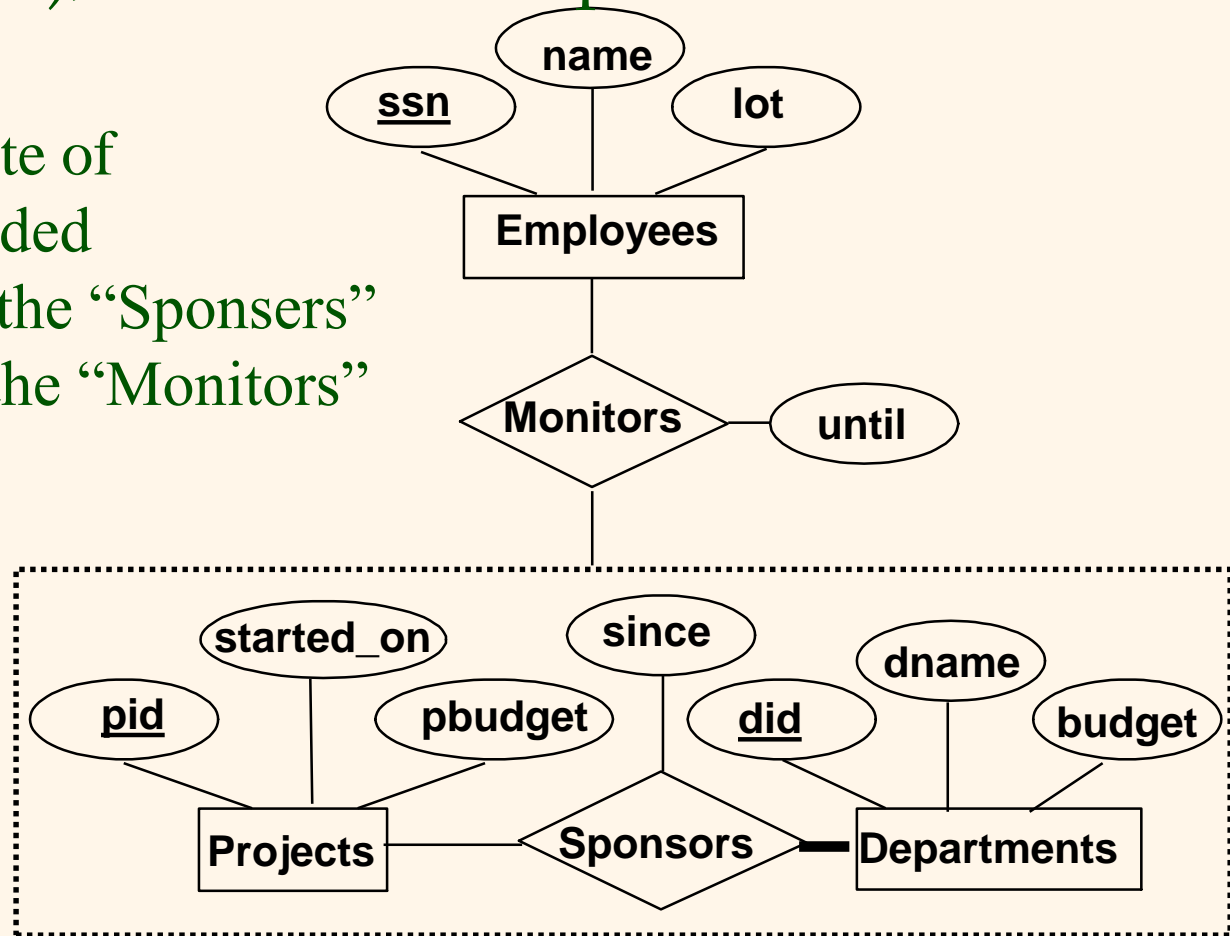
- *Hourly_Emps*: *ssn*, *name*, *lot*, *hourly_wages*, *hours_worked*.
- Each employee must be in one of these two subclasses.

Aggregation




❖ *Monitors*: Create a relation with the key attribute of “Employee” (ssn) and the key attribute for “Sponsors” (did,pid), and the descriptive attribute of “Monitors” (until)

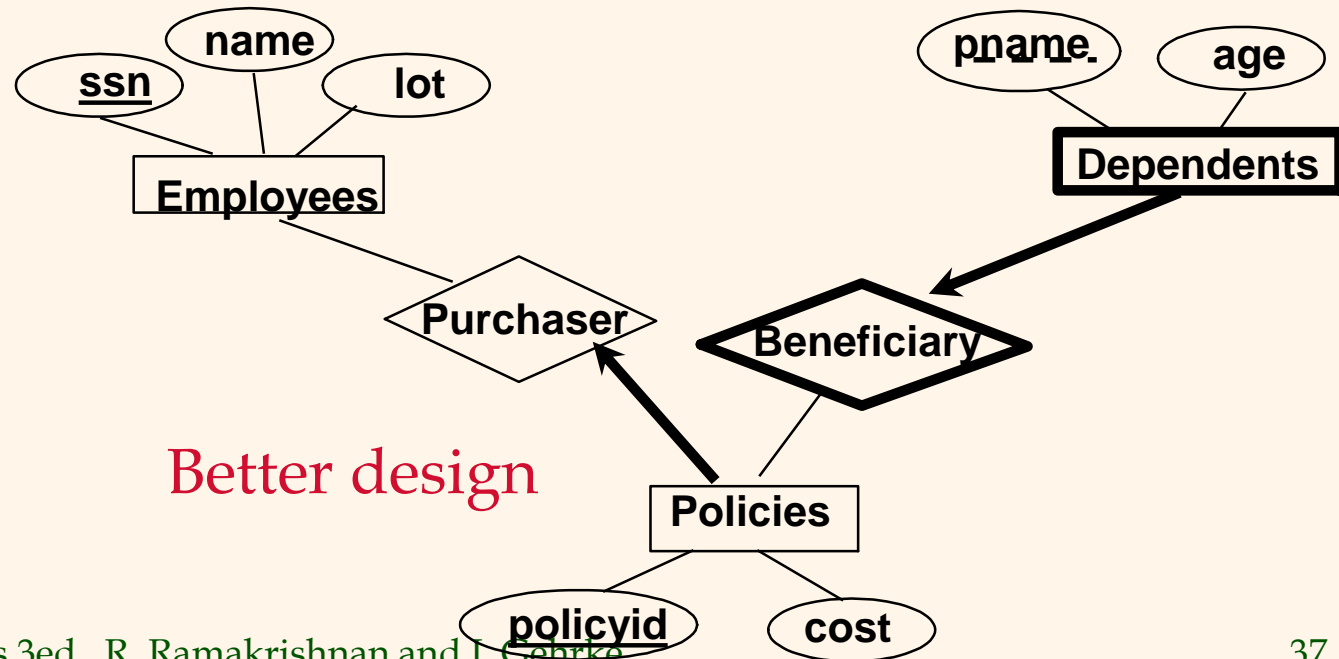
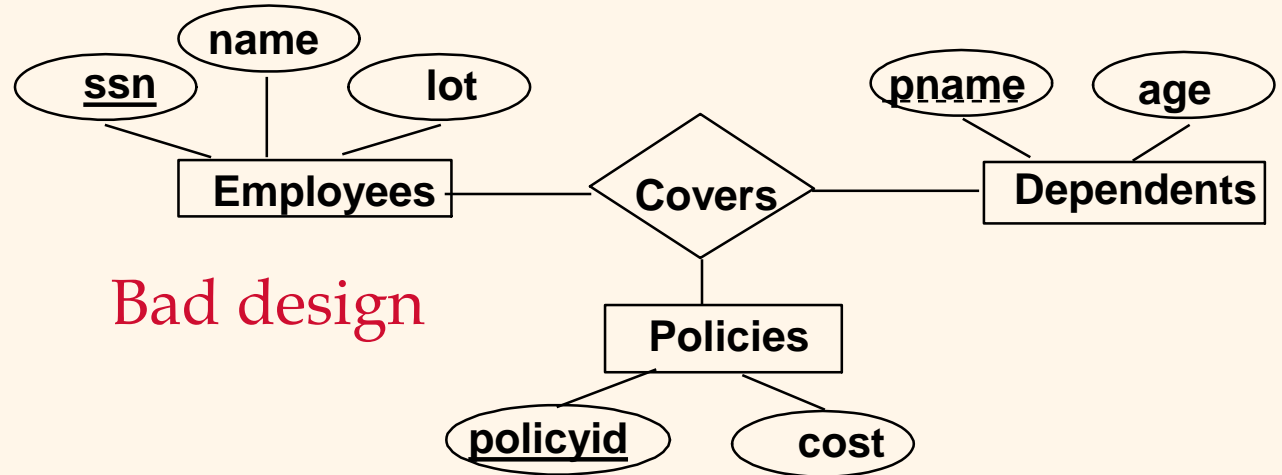
- The descriptive attribute of “sponsors” is not included
- Not every (pid,did) in the “Sponsors” relation participate in the “Monitors” relation
- If “Sponsors” has no descriptive attributes and has total participation in “Monitors”, we can drop “Sponsors”




Review: Binary vs. Ternary Relationships



- ❖ If each policy is owned by just one employee, and each dependent is tied to the covering policy, first diagram is inaccurate.



Binary vs. Ternary Relationships (Contd.)




- ❖ The key constraints allow us to combine Purchaser with Policies and Beneficiary with Dependents.

```
CREATE TABLE Policies (  
    policyid INTEGER,  
    cost REAL,  
    ssn CHAR(11) NOT NULL,  
    PRIMARY KEY (policyid),  
    FOREIGN KEY (ssn) REFERENCES Employees,  
    ON DELETE CASCADE)
```

- ❖ Participation constraints lead to NOT NULL constraints.

```
CREATE TABLE Dependents (  
    pname CHAR(20),  
    age INTEGER,  
    policyid INTEGER,  
    PRIMARY KEY (pname, policyid),  
    FOREIGN KEY (policyid) REFERENCES Policies,  
    ON DELETE CASCADE)
```

Binary vs. Ternary Relationships (Contd.)




- ❖ What if Policies is a weak entity set?
 - Cascaded weak entities

```
CREATE TABLE Policies (  
    policyid INTEGER,  
    cost REAL,  
    ssn CHAR(11),  
    PRIMARY KEY (policyid,ssn).  
    FOREIGN KEY (ssn) REFERENCES  
    Employees ON DELETE CASCADE)
```

```
CREATE TABLE Dependents (  
    pname CHAR(20),  
    ssn CHAR(11),  
    age INTEGER,  
    policyid INTEGER,  
    PRIMARY KEY (pname, policyid,ssn).  
    FOREIGN KEY (policyid,ssn) REFERENCES
```

Policies




Views

- ❖ A view is just a relation, but we store a *definition*, rather than a set of tuples.

```
CREATE VIEW YoungActiveStudents (name, grade)
AS SELECT S.name, E.grade
FROM Students S, Enrolled E
WHERE S.sid = E.sid and S.age < 21
```

- ❖ Views can be dropped using the **DROP VIEW** command.
- How to handle **DROP TABLE** if there's a view on the table?
 - DROP TABLE command has options to let the user specify this.




Updates on Views

- ❖ SQL-92 allows updates to be specified only on views defined on a single table

```
CREATE VIEW GoodStudents (sid, gpa)  
    AS SELECT  sid,gpa  
    FROM Students  
    WHERE  gpa > 3.0
```

- ❖ Deleting a tuple from GoodStudent would delete the corresponding tuple from the original table
- ❖ What would happen if “sid” is replaced by “sname” in the view




Updates on Views

- ❖ A tuple can be inserted into the view

```
INSERT INTO GoodStudents (sid, gpa)  
VALUES (34261,3.5)
```

- ❖ What would happen to the other fields → NULL
- ❖ Can we insert a GPA less than 3 → YES
- ❖ To prevent this → Add “WITH CHECK OPTION” to the view definition



Relational Model: Summary

- ❖ A tabular representation of data.
- ❖ Simple and intuitive, currently the most widely used.
- ❖ Integrity constraints can be specified by the DBA, based on application semantics. DBMS checks for violations.
 - Two important ICs: primary and foreign keys
 - In addition, we *always* have domain constraints.
- ❖ Powerful and natural query languages exist.
- ❖ Rules to translate ER to relational model