

# *Overview of Query Evaluation*

## Chapter 12


# *Outline*

- ❖ Query Optimization Overview
- ❖ Algorithm for Relational Operations

# *Overview of Query Evaluation*

- ❖ DBMS keeps descriptive data in system catalogs.
- ❖ SQL queries are translated into an extended form of relational algebra:
  - Query Plan Reasoning:
    - Tree of operators
    - with choice of one among several algorithms for each operator

# Query Plan Evaluation



## ■ Query Plan Execution:

- Each operator typically implemented using a 'pull' interface
- when an operator is 'pulled' for next output tuples, it 'pulls' on its inputs and computes them.

# Overview of Query Evaluation

## ❖ Query Plan Optimization :

- **Ideally:** Want to find best plan.  
**Practically:** Avoid worst plans!

## ❖ Two main issues in query optimization:

- For a given query, **what plans are considered?**
  - Algorithm to search plan space for cheapest (estimated) plan.
- How is the **cost of a plan estimated?**
  - Cost models based on I/O estimates

# Query Processing

## ❖ Common Techniques for Query Processing Algorithms:

- **Indexing:** Can use WHERE conditions to retrieve small set of tuples from large relation
- **Iteration:** Examine all tuples in an input table, one after the other (like in sorting algorithm).
- **Partitioning:** By using sorting or hashing, we partition input tuples and replace expensive operation by similar operations on smaller inputs.

*\* Watch for these techniques as we discuss query evaluation!*

# *Access Paths*

## ❖ An access path

- A method of retrieving tuples from a table
- Method:
  - File scan,
  - Index that matches a selection (in the query)
- Note :
  - Contributes significantly to cost of relational operator.

# *Matching an Access Path*

- ❖ A **tree index matches** (a conjunction of) terms that involve only attributes in a **prefix** of search key.
- Example : Given tree index on  $\langle a, b, c \rangle$ 
  - selection  $a=5 \text{ AND } b=3$  ?
  - selection  $a=5 \text{ AND } b>6$  ?
  - selection  $b=3$  ?



# Matching an Access Path

- ❖ A **hash index** matches (a conjunction of) terms that has a term *attribute = value* for every attribute in search key of index.
- Example : Given hash index on  $\langle a, b, c \rangle$ 
  - selection  $a=5$  AND  $b=3$  and  $c=5$  ?
  - selection  $c=5$  AND  $b=6$  ?
  - selection  $a=5$  ?
  - selection  $a>5$  AND  $b=3$  and  $c=5$  ?

# *Query Evaluation of Selection*

# *Selection*

- ❖ Example :  $\sigma_{R.attr \text{ OP value }}(R)$
- ❖ **Case 1:** No Index, NOT sorted on R.attr
- ❖ Must scan the entire relation.  
Most selective access path = file scan  
Cost: M

# *Selection*

- ❖ **Case 2:** No Index, Sorted Data on R.attr
- ❖ Binary search for first tuple.  
Scan R for all satisfied tuples.  
Cost:  $O(\log_2 M)$

# *Selection Using B+ tree index*

## ❖ Case 3: B+ tree Index

- Cost I (finding qualifying data entries)  
+ cost II (retrieving records) :
  - Cost I: 2-3 I/Os. (depth of B+ tree)
  - Cost II:
    - clustered index 1 I/O,
    - unclustered index upto one I/O per qualifying tuple.

# *Example : Using B+ Index for Selections*

```
SELECT *  
FROM   Reserves R  
WHERE  R.rname < 'C%'
```

## ❖ Example :

- Assume uniform distribution of names, about 10% of tuples qualify (100 pages, 10,000 tuples).
- Clustered index:
  - little more than 100 I/Os;
- Unclustered index :
  - up to 10,000 I/Os!

# *Selection --- B+ Index*

## ❖ *Refinement for unclustered indexes:*

1. Find qualifying data entries.
2. Sort rid's of data records to be retrieved.
3. Fetch rids in order.  
    Avoid retrieving the same page multiple times.

However, # of such pages likely to be still higher than with clustering.

- ## ❖ Use of unclustered index for a range selection could be expensive. *Simpler if just scan data file.*

# *Selection – Hash Index*

- ❖ Hash index is good for equality selection.
- ❖ Cost:
  - Cost I (retrieve index bucket page)
  - + Cost II (retrieving qualifying tuples from R)
    - Cost I is one I/O
    - Cost II could up to one I/O per satisfying tuple.



# *General Condition : Conjunction*

- ❖ A condition with several predicates combined by conjunction (AND):
- ❖ Example : *day<8/9/94 AND bid=5 AND sid=3.*

# General Selections (Conjunction)

First approach: (utilizing single index)

- ❖ Find the *most selective access path*, retrieve tuples using it.
  - To *reduce* the number of tuples *retrieved*
- ❖ Apply any remaining terms that don't match the index:
  - To *discard* some retrieved tuples
  - This does not affect number of tuples/pages fetched.
- ❖ Example : Consider *day<8/9/94 AND bid=5 AND sid=3*.
  - A B+ tree index on *day* can be used;
  - then *bid=5* and *sid=3* must be checked for each retrieved tuple.
  - Hash index on *<bid, sid>* could be used
  - *day<8/9/94* must then be checked on fly.

# General Selections

## Second approach (utilizing multiple index)

- ❖ Assuming 2 or more matching indexes that use Alternatives (2) or (3) for data entries.
  - Get sets of rids of data records using each matching index.
  - Then *intersect* these sets of rids
  - Retrieve records and apply any remaining terms.
  
- ❖ Example : Consider *day<8/9/94 AND bid=5 AND sid=3*.
  - A B+ tree index I on *day* and an index II on *sid*, both Alternative (2).
  - - Retrieve rids of records satisfying *day<8/9/94* using index I,
  - - Retrieve rids of recs satisfying *sid=3* using Index II
  - - Intersect rids
  - - Retrieve records and check *bid=5*.

# *General Condition : Disjunction*

- ❖ Disjunction condition: one or more terms (R.attr op value) connected by OR (  $\vee$  ).
- ❖ Example : (*day*<8/9/94) OR (*bid*=5 AND *sid*=3)

# *General Selection (Disjunction)*

- ❖ **Case 1:** Index is not available for one of terms.  
Need a file scan. Check other conditions in this file scan.
- ❖ E.g., Consider *day < 8/9/94 OR rname = 'Joe'*
  - No index on day. Need a File scan.
  - Even index is available in rname, does not help.

# *General Selection (Disjunction)*

- ❖ **Case 2:** Every term has a matching index.
  - Retrieve candidate tuples using index.
  - Then Union the results
  
- ❖ **Example :** consider *day < 8/9/94 OR rname = 'Joe'*
  - Assume two B+ tree indexes on *day* and *rname*.
  - Retrieve tuples satisfying *day < 8/9/94*
  - Retrieve tuples satisfying *rname = 'Joe'*
  - Union the retrieved tuples.

# *Query Evaluation of Projection*

# Algorithms for Projection

```
SELECT  DISTINCT  
        R.sid, R.bid  
FROM    Reserves R
```

- ❖ The expensive part is removing duplicates.
  - SQL systems don't remove duplicates unless keyword DISTINCT is specified in query.
- ❖ Sorting Approach:
  - Sort on <sid, bid> and remove duplicates.  
(Can optimize this by dropping unwanted information while sorting.)
- ❖ Hashing Approach:
  - Hash on <sid, bid> to create partitions.  
Load partitions into memory one at a time, build in-memory hash structure, and eliminate duplicates.
- ❖ Indexing Approach :
  - If there is an index with both R.sid and R.bid in the search key, may be cheaper to sort data entries!



# *Query Evaluation of Joins*

# Schema for Examples

Sailors (*sid*: integer, *sname*: string, *rating*: integer, *age*: real)

Reserves (*sid*: integer, *bid*: integer, *day*: dates, *rname*: string)

- ❖ Similar to old schema; *rname* added for variations.
- ❖ Reserves:
  - Each tuple is 40 bytes long, 100 tuples per page, 1000 pages.
- ❖ Sailors:
  - Each tuple is 50 bytes long, 80 tuples per page, 500 pages.

# Equality Joins With One Join Column

```
SELECT *  
FROM   Reserves R1, Sailors S1  
WHERE  R1.sid=S1.sid
```

- ❖ In algebra:  $R \bowtie S$ . Common! Must be carefully optimized.
- ❖  $R \times S$  is large; so  $R \times S$  followed by a selection is inefficient.
- ❖ Assume:
  - M pages of R,  $p_R$  tuples per page (i.e., number of tuples of R =  $M * p_R$ ),  
N pages of S,  $p_S$  tuples per page (i.e., number of tuples of S =  $N * p_S$ ),
  - In our examples, R is Reserves and S is Sailors.
- ❖ *Cost metric*: # of I/Os. We will ignore output costs.

# *Typical Choices for Joins*


## ❖ Nested Loops Join

- Simple Nested Loops Join: Tuple-oriented
- Simple Nested Loops Join: Page-oriented
- Block Nested Loops Join
- Index Nested Loops Join

## ❖ Sort Merge Join

## ❖ Hash Join

# Simple Nested Loops Join



```
foreach tuple r in R do
    foreach tuple s in S do
        if  $r_i == s_j$  then add  $\langle r, s \rangle$  to result
```


## ❖ Algorithm :

For each **tuple** in **outer relation R**, we scan **inner relation S**.

## ❖ Cost :

- Scan of outer + for each tuple of outer, scan of inner relation.
- $\text{Cost} = M + p_R * M * N$
- $\text{Cost} = 1000 + 100 * 1000 * 500 \text{ I/Os.}$

# Simple Nested Loops Join



```
foreach tuple r in R do
    foreach tuple s in S do
        if  $r_i == s_j$  then add  $\langle r, s \rangle$  to result
```

## ❖ Tuple-oriented:

For each *tuple* in *outer* relation R, we scan *inner* relation S.

- Cost:  $M + p_R * M * N = 1000 + 100 * 1000 * 500$  I/Os.


## ❖ Page-oriented:

For each *page* of R, get each *page* of S, and write out matching pairs of tuples  $\langle r, s \rangle$ , where r is in R-page and S is in S-page.

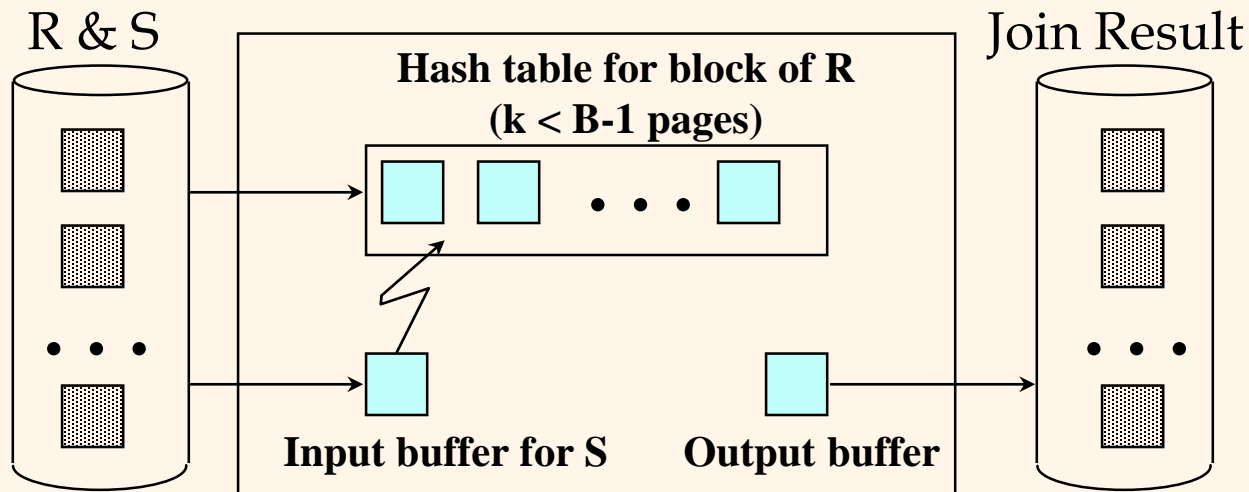
## ❖ Cost :

- Scan of outer pages + for each page of outer, scan of inner relation.
- Cost =  $M + M * N$
- Cost =  $1000 + 1000 * 500$  IOs.
- smaller relation (S) is outer, cost =  $500 + 500 * 1000$  IOs.

# Block Nested Loops Join



- ❖ One page as input buffer for scanning inner S,
- ❖ One page as the output buffer,
- ❖ Remaining pages to hold ``block'' of outer R.
  - For each matching tuple  $r$  in R-block,  $s$  in S-page, add  $\langle r, s \rangle$  to result.
  - Then read next R-block, scan S again. Etc.
  - Find matching tuple ?  $\rightarrow$  Use in-memory hashing.



# *Cost of Block Nested Loops*

❖ Cost: Scan of outer + #outer blocks \* scan of inner

▪ #outer blocks =

$$\lceil \# \text{ of pages of outer} / \text{blocksize} \rceil$$



# Examples of Block Nested Loops

- ❖ Cost: Scan of outer + #outer blocks \* scan of inner
- ❖ With Reserves (R) as outer, & 100 pages of R as block:
  - Cost of scanning R is 1000 I/Os; a total of 10 *blocks*.
  - Per block of R, we scan Sailors (S); 10\*500 I/Os.
  - E.g., If a block is 90 pages of R, we would scan S 12 times.
- ❖ With 100-page block of Sailors as outer:
  - Cost of scanning S is 500 I/Os; a total of 5 blocks.
  - Per block of S, we scan Reserves; 5\*1000 I/Os.

# *Examples of Block Nested Loops*

## ❖ Optimizations?

- With sequential reads considered, analysis changes: may be best to divide buffers evenly between R and S.
- Double buffering would also be suitable.

# *Index Nested Loops Join*

```
foreach tuple r in R do  
    foreach tuple s in S where  $r_i == s_j$  do  
        add <r, s> to result
```

- ❖ An index on join column of one relation (say S), use S as inner and exploit the index.
- ❖ Cost:
  - Scan the outer relation R
  - For each R tuple, sum cost of finding matching S tuples
  - Cost:  $M + (M * p_R) * \text{cost of finding matching S tuples}$

# *Index Nested Loops Join*

- ❖ For each R tuple, cost of probing S index is :
  - about 1.2 for hash index,
  - 2-4 for B+ tree.
- ❖ Cost of retrieving S tuples (assuming Alt. (2) or (3) for data entries) depends on clustering:
  - Clustered : 1 I/O (typical),
  - Unclustered: up to 1 I/O per matching S tuple.

# Examples of Index Nested Loops

## ❖ Hash-index (Alt. 2) on *sid* of Sailors (as inner):

- Scan Reserves:
  - 1000 page I/Os,
  - 100\*1000 tuples.
- For each Reserves tuple:
  - 1.2 I/Os to get data entry in index,
  - plus 1 I/O to get (the exactly one) matching Sailors tuple.
  - Total:  $100,000 * (1.2 + 1) = 220,000$  I/Os.
- In total, we have:
  - 1000 I/Os plus
  - 220,000 I/Os.
  - Equals 221,000 I/Os

# Examples of Index Nested Loops

## ❖ Hash-index (Alt. 2) on *sid* of Reserves (as inner):

- Scan Sailors:
    - 500 page I/Os,
    - 80\*500 tuples.
  - For each Sailors tuple:
    - 1.2 I/Os to find index page with data entries,
    - plus cost of retrieving matching Reserves tuples.
  - Assuming uniform distribution:  
2.5 reservations per sailor (100,000 / 40,000).
  - Cost of retrieving them is 1 or 2.5 I/Os  
depending on whether the index is clustered.
- Total :  $500 + 40,000 * (1.2 + 2.5)$ .

# *Simple vs. Index Nested Loops Join*


- ❖ Assume: M Pages in R,  $p_R$  tuples per page,  
N Pages in S,  $p_S$  tuples per page,  
B Buffer Pages.
- ❖ Nested Loops Join
  - Simple Nested Loops Join
    - Tuple-oriented:  $M + p_R * M * N$
    - Page-oriented:  $M + M * N$
    - Smaller as outer helps.
  - Block Nested Loops Join
    - $M + N * \lceil M / (B-2) \rceil$
    - Dividing buffer evenly between R and S helps.
  - Index Nested Loops Join
    - $M + (M * p_R) * \text{cost of finding matching S tuples}$
    - $\text{cost of finding matching S tuples} = \text{cost of Probe} + \text{cost of retrieving}$
- ❖ With unclustered index, if number of matching inner tuples for each outer tuple is small, cost of INLJ is much smaller than SNLJ.

# *Join: Sort-Merge ( $R \bowtie_{i=j} S$ )*



- (1). Sort R and S on the join column.
- (2). Scan R and S to do a ``merge'' on join column
- (3). Output result tuples.



# Example of Sort-Merge Join



<u>sid</u>	sname	rating	age
22	dustin	7	45.0
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0



<u>sid</u>	<u>bid</u>	<u>day</u>	rname
28	103	12/4/96	guppy
28	103	11/3/96	yuppy
31	101	10/10/96	dustin
31	102	10/12/96	lubber
31	101	10/11/96	lubber
58	103	11/12/96	dustin

# Join: Sort-Merge ( $R \bowtie_{i=j} S$ )

- (1). Sort R and S on the join column.
- (2). Scan R and S to do a “merge” on join col.
- (3). Output result tuples.

## ■ Merge on Join Column:


- Advance scan of R until current R-tuple  $\geq$  current S tuple,
- then advance scan of S until current S-tuple  $\geq$  current R tuple;
- do this until current R tuple = current S tuple.
  
- At this point, all R tuples with same value in  $R_i$  (*current R group*) and all S tuples with same value in  $S_j$  (*current S group*) match;
- So output  $\langle r, s \rangle$  for all pairs of such tuples.
  
- Then resume scanning R and S (as above)

# *Join: Sort-Merge ( $R \bowtie_{i=j} S$ )*


## ❖ Note :

- R is scanned once; each S group is scanned once per matching R tuple.
- Multiple scans of an S group are likely to find needed pages in buffer.

# Cost of Sort-Merge Join



<u>sid</u>	sname	rating	age
22	dustin	7	45.0
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0



<u>sid</u>	<u>bid</u>	<u>day</u>	rname
28	103	12/4/96	guppy
28	103	11/3/96	yuppy
31	101	10/10/96	dustin
31	102	10/12/96	lubber
31	101	10/11/96	lubber
58	103	11/12/96	dustin

Cost of sort-merge :

- ❖ Sort R
- ❖ Sort S
- ❖ Merge R and S

# Example of Sort-Merge Join


<u>sid</u>	sname	rating	age
22	dustin	7	45.0
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

<u>sid</u>	<u>bid</u>	<u>day</u>	rname
28	103	12/4/96	guppy
28	103	11/3/96	yuppy
31	101	10/10/96	dustin
31	102	10/12/96	lubber
31	101	10/11/96	lubber
58	103	11/12/96	dustin

- ❖ Best case: ?
- ❖ Worst case: ?
- ❖ Average case ?

# Cost of Sort-Merge Join

- ❖ Best Case Cost:  $(M+N)$ 
  - Already sorted.
  - The cost of scanning,  $M+N$




- ❖ Worst Case Cost:  $M \log M + N \log N + (M+N)$
- ❖ Many pages in R in same partition. (Worst, all of them).  
The pages for this partition in S don't fit into RAM.  
Re-scan S is needed. Multiple scan S is expensive!
- ❖ Note: Guarantee  $M+N$  if key-FK join, or no duplicates.

# Cost of Sort-Merge Join

## ❖ Average Cost:

- ~ In practice, roughly linear in M and N
- So  $O(M \log M + N \log N + (M+N))$



# *Comparison with Sort-Merge Join*

- ❖ **Average Cost:**  $O(M \log M + N \log N + (M+N))$
- ❖ Assume  $B = \{35, 100, 300\}$ ; and  
 $R = 1000$  pages,  $S = 500$  pages
- ❖ Sort-Merge Join
- ❖ both  $R$  and  $S$  can be sorted in 2 passes,
- ❖  $\log M = \log N = 2$
- ❖ total join cost:  $2*2*1000 + 2*2*500 + (1000 + 500)$   
 $= 7500$ .
- ❖ Block Nested Loops Join:  $2500 \sim 15000$



# *Refinement of Sort-Merge Join*

## ❖ IDEA :


- Combine the merging phases when *sorting* R ( or S) with the merging in join algorithm.

# Refinement of Sort-Merge Join

- ❖ **IDEA :** Combine the merging phases when *sorting* R ( or S) with the merging in join algorithm.
  - If we do the following: perform Pass 0 of sort on R; perform Pass 0 of sort on S; merge and join on the fly – the total IO cost for join is  $3 (M + N)$
  - When is the above possible? When  $M/B + N/B + 1 \leq B$ ; In other words when  $B (B - 1) \geq (M + N)$   
(The above expression is modified from that in the book)
  - **Cost:  $3 (M + N)$  as follows**
    - (read+write R and S in Pass 0)
    - + (read R and S in merging pass and join on fly)
    - + (writing of result tuples).
  - In example, cost goes down from 7500 to 4500 I/Os.

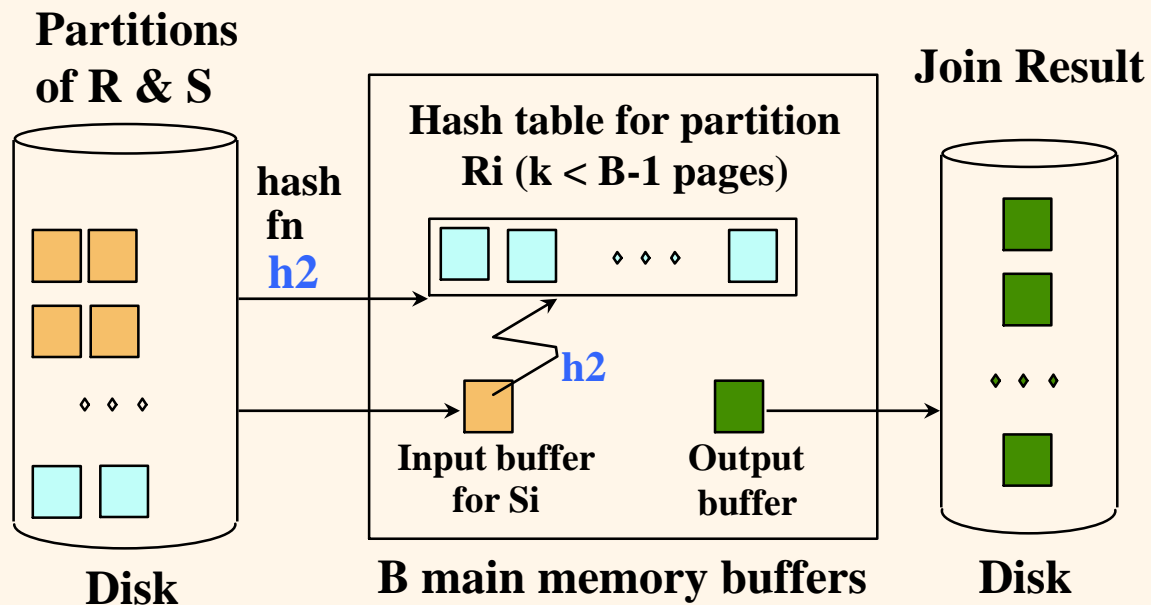
# Hash-Join

Partition both  
relations using same  
hash fn **h**:  
R tuples in partition i  
will only match S  
tuples in partition i.



# Hash-Join

- ❖ Read in a partition of R, hash it using  $h_2$  ( $\neq h_1$ ). Scan matching partition of S, search for matches.



# *Cost of Hash-Join*

- ❖ In partitioning phase, read+write both relations:
  - $2(M+N)$ .
- ❖ In matching phase, read both relations:
  - $M+N$ .
- ❖ Total :  $3(M+N)$
- ❖ E.g., total of 4500 I/Os in our running example.

# Observation on Hash-Join

- ❖ Memory Requirement: When is total cost  $3(M + N)$ ?
  - Partition fit into available memory?
  - Assuming  $B$  buffer pages. #partitions  $k \leq B-1$  (why?), (to min size of each partition, we choose #partitions =  $B - 1$ )
  - Assuming uniformly sized partitions, and maximizing  $k$ , we get:
    - $k = B-1$ , and size of partition =  $M/(B-1)$  ( $M$  is the number of pages of  $R$ )
    - in-memory hash table to speed up the matching of tuples, a little more memory is needed:  $f * M/(B-1)$  (You can assume  $f = 1$ , unless explicitly specified)
    - $f$  is fudge factor used to capture the small increase in size between the partition and a hash table for partition.
  - Probing phase, one for  $S$ , one for output,  $B \geq f * M/(B-1) + 2$  for hash join to perform well (i.e., cost of hash join =  $3(M + N)$ ). In other words,  $(B - 1)(B - 2) \geq f * M$

# *Observation on Hash Join*

## ❖ Overflow

- If the hash function does **not partition uniformly**, one or more R partitions may not fit in memory.
- Significantly degrade the performance.
- Can apply hash-join technique recursively to do the join of this overflow R-partition with corresponding S-partition.

# Hybrid Hash-Join

- ❖ Idea: Do not write one of the partitions of R and S to disk.
- ❖ When is it possible? We can keep one of the partitions of the smaller relation always in memory.
- ❖  $B \geq f * M/k$  (buffers for keeping a partition)
  - +  $(k - 1)$  (keep 1 page in buffer for each of the remaining partitions)
  - + 1 (1 page in buffer for reading in S (or later R))
  - + 1 (1 output page when reading in R)

Remember:  $k$  = number of partitions  
i.e.,  $(B - (k + 1)) \geq f * M/k$
- ❖ Choose such an appropriate  $k$  (or number of partitions)



# Hybrid Hash-Join (contd)

## ❖ How to perform Hybrid Hash-Join?

### ▪ Partitioning S is done as:

- Build an in-memory hash table for the first partition of S during the partitioning phase.
- Other partitions keep 1 page in buffer and write to disk when needed.
- 1 buffer page for reading in S

### ▪ Partitioning R is done as:

- If a tuple hashes to the partition corresponding to the in-memory partition of S, then join and output tuples
- If a tuple hashes to any of the remaining  $(k - 1)$  partitions, write it to the buffer page (and write this buffer page to disk as needed)
- 1 buffer page for reading in R; 1 buffer page for output

### ▪ Remaining partitions of R and S are done as usual

## ❖ **Saving:** avoid writing the first partitions of R and S to disk.

- E.g.  $R = 500$  pages,  $S = 1000$  pages  $B = 300$  (We make 2 partitions)  
partition phase: scan R and write one partition out.  $500 + 250$   
scan S and write out one partition.  $1000 + 500$   
probing phase: only second partition is scanned:  $250 + 500$
- Total = 3000 ( Hash Join will take 4500 )

# *Hash-Join vs. Sort-Merge Join*

## ❖ Sort-Merge Join vs. Hash Join:

- Given a certain amount of memory:  $B - 1 \geq (M + N)$  both have a cost of  $3(M+N)$  I/Os.
- If partition is not uniformly sized (data skew); Hash-Join less sensitive.
- Hash Join superior if relation sizes differ greatly;  $B$  is between  $\sqrt{N}$  and  $\sqrt{L}$  (roughly), where  $L = (M + N)$

# General Join Conditions

## ❖ Equalities over several attributes

- (e.g., *R.sid=S.sid AND R.rname=S.sname*):
- **INL-Join** : build index on *<sid, sname>* (if S is inner); or use existing indexes on *sid* or *sname*.
- **SM-Join** and **H-Join** : sort/partition on combination of the two join columns.

## ❖ Inequality conditions

- (e.g., *R.rname < S.sname*):
- **INL-Join**: need (clustered!) B+ tree index.
  - Range probes on inner; # matches likely to be much higher than for equality joins.
- Hash Join, Sort Merge Join not applicable.
- Block NL quite likely to be the best join method here.

# *Summary*

- ❖ There are several alternative evaluation algorithms for each relational operator.

# *Conclusion*

Not one method wins !

Optimizer must assess situation to  
select best possible candidate